

# CS 1674/2074: Grouping: segmentation

---

**PhD. Nils Murrugarra-Llerena**  
[nem177@pitt.edu](mailto:nem177@pitt.edu)



# [Motivation] Grouping

## Finding a Crowd at a Concert

Imagine you're at a large outdoor music festival and you need to find your friends.

What would you do?

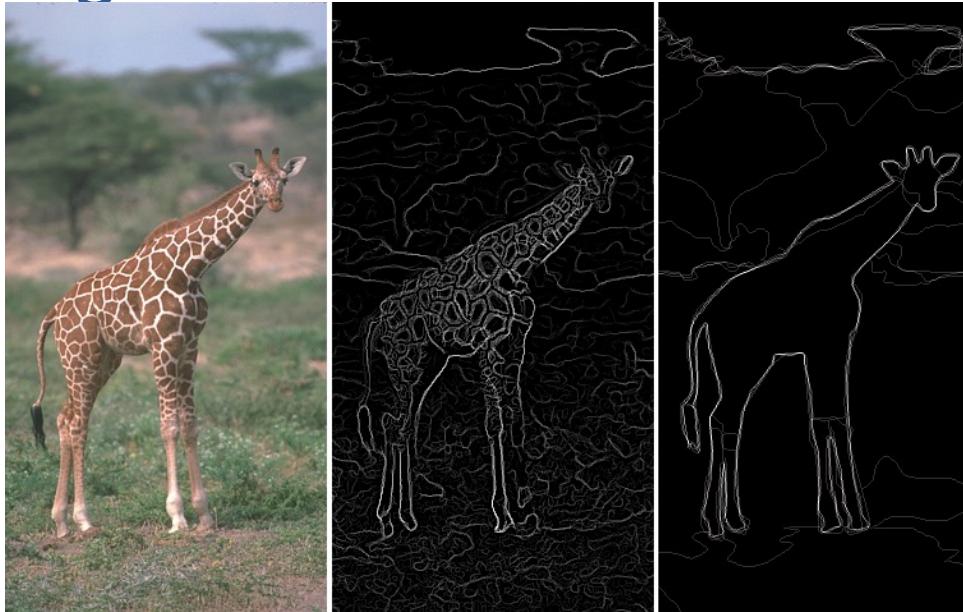
Your eyes don't process the scene one person at a time. Instead, you automatically group people together. You might spot a cluster of people wearing the *same band t-shirt*, or a *tight-knit group all dancing together in a circle*. You instantly recognize these as distinct groups, not just random individuals.



## Plan for today

- Segments
  - Find which pixels form a consistent region
  - Clustering (e.g. K-means)

# Edges vs Segments

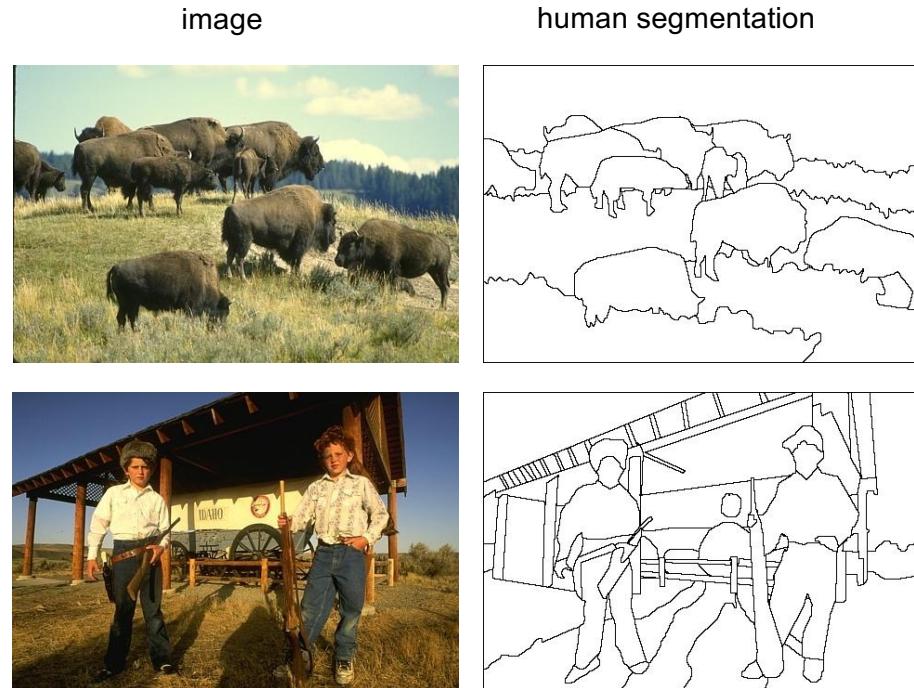


- **Edges:** More low-level; don't need to be closed
- **Segments:** Ideally one segment for each semantic group/object; should include closed contours

Figure adapted from J. Hays

# The goals of segmentation

- Separate image into coherent “objects”

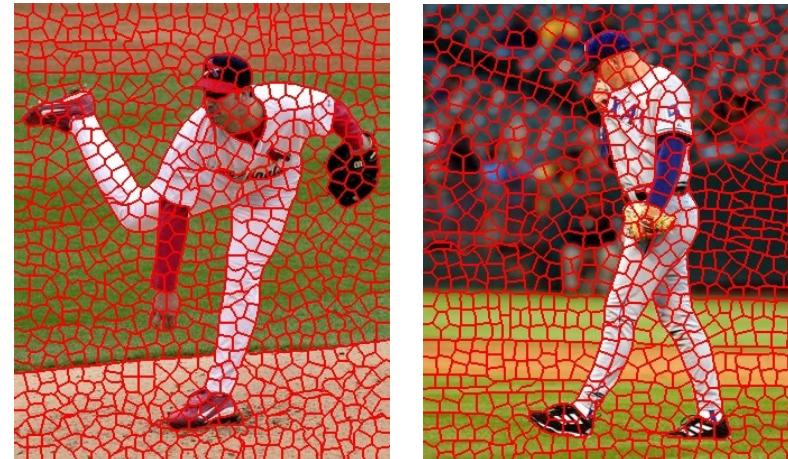


Source: L. Lazebnik

# The goals of segmentation

- Separate image into coherent “objects”
- Group together similar-looking pixels for efficiency of further processing

“superpixels”



X. Ren and J. Malik. [Learning a classification model for segmentation](#). ICCV 2003.

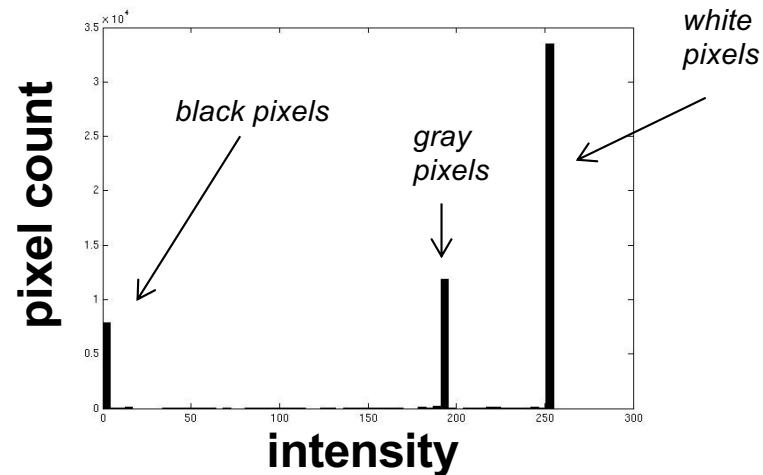
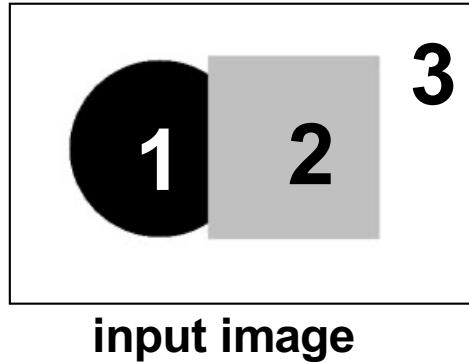
Source: L. Lazebnik

# Similarity



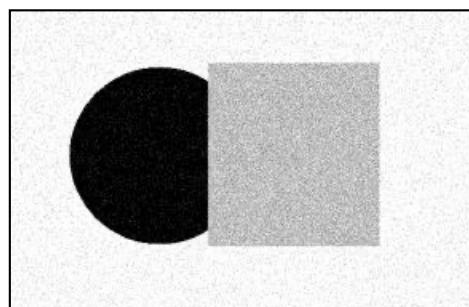
Slide: K. Grauman

# Image Segmentation: Toy Example

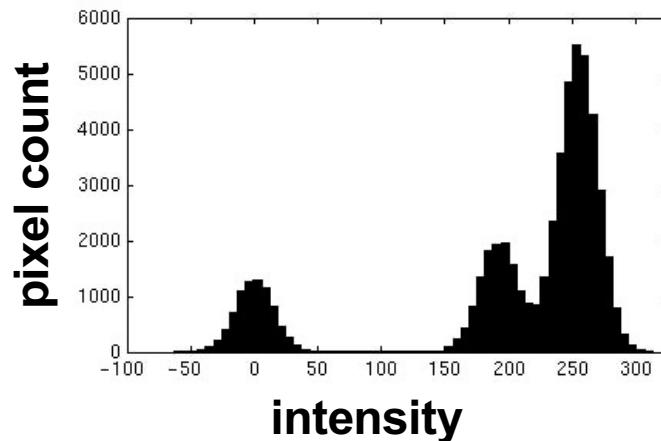


- These intensities define the three groups.
- We could label every pixel in the image according to which of these primary intensities it is.
  - i.e., *segment* the image based on the intensity feature.
- What if the image isn't quite so simple?

# Image Segmentation: Toy Example



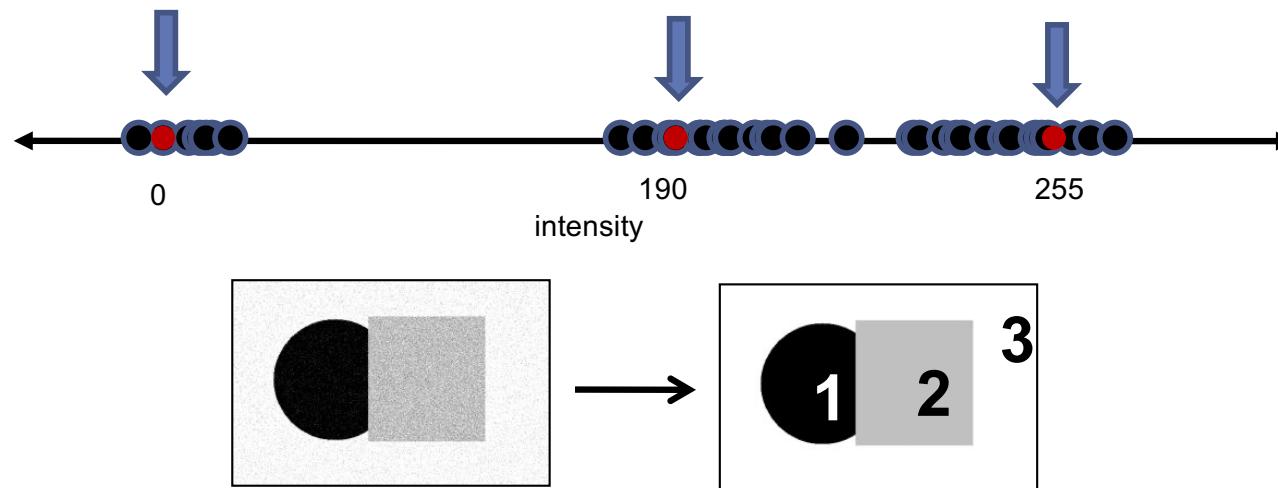
input image



- Now how to determine the three main intensities that define our groups?
- We need to **cluster**.

Source: K. Grauman

# Image Segmentation: Toy Example



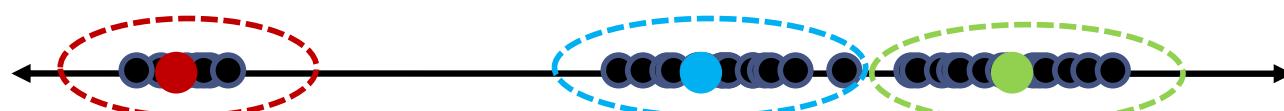
- Goal: choose three “centers” as the **representative** intensities, and label every pixel according to which of these centers it is nearest to.
- Best cluster centers are those that **minimize sum of squared differences** (SSD) between all points and their nearest cluster center  $c_i$ :

$$\sum_{\text{clusters } i} \sum_{\text{points } p \text{ in cluster } i} \|p - c_i\|^2$$

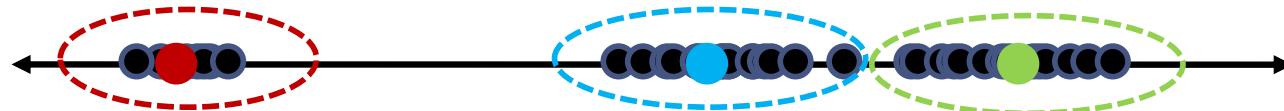
Source: K. Grauman

# Clustering

- With this objective, it is a “chicken and egg” problem:
  - If we knew the **cluster centers**, we could allocate points to groups by assigning each to its closest center.



- If we knew the **group memberships**, we could get the centers by computing the mean per group.



# K-means clustering

- **Basic idea:** randomly initialize the  $k$  cluster centers, and iterate between the two steps we just saw.
  1. Randomly initialize the **cluster centers**,  $c_1, \dots, c_K$
  2. Given **cluster centers**, determine **points** in each cluster
    - For each point  $p$ , find the closest  $c_i$ . Put  $p$  into **cluster i**
  3. Given **points** in each **cluster**, solve for  $c_i$ 
    - Set  $c_i$  to be the mean of **points** in **cluster i**
  4. If  $c_i$  have changed, repeat Step 2

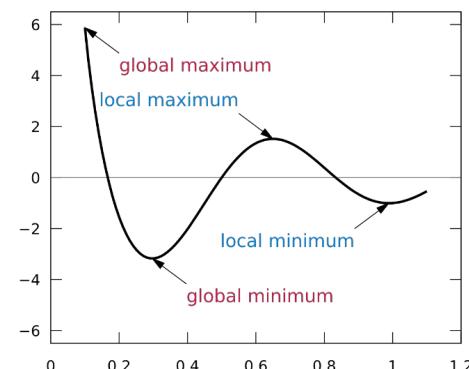


## Properties

- Will always converge to some solution
- Can be a “local minimum” of objective:

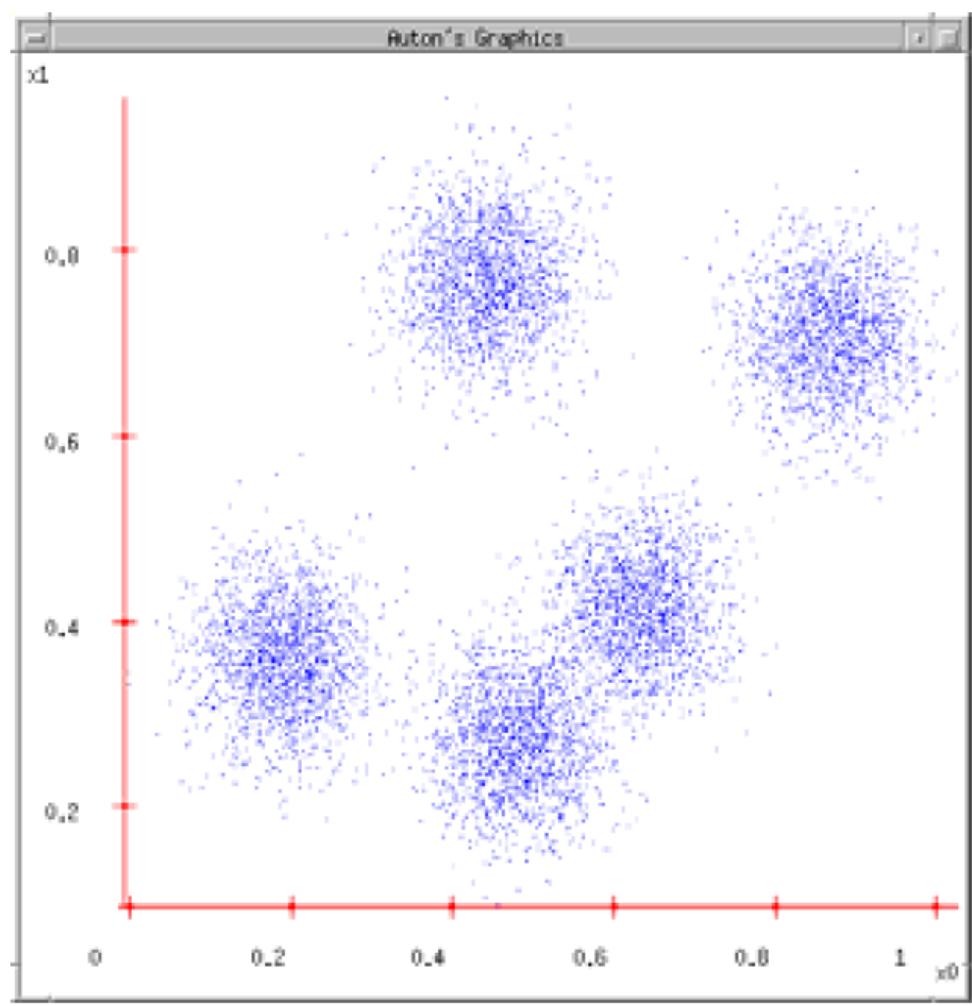
$$\sum_{\text{clusters } i} \sum_{\text{points } p \text{ in cluster } i} \|p - c_i\|^2$$

Slide: Steve Seitz, image: Wikipedia



# K-means

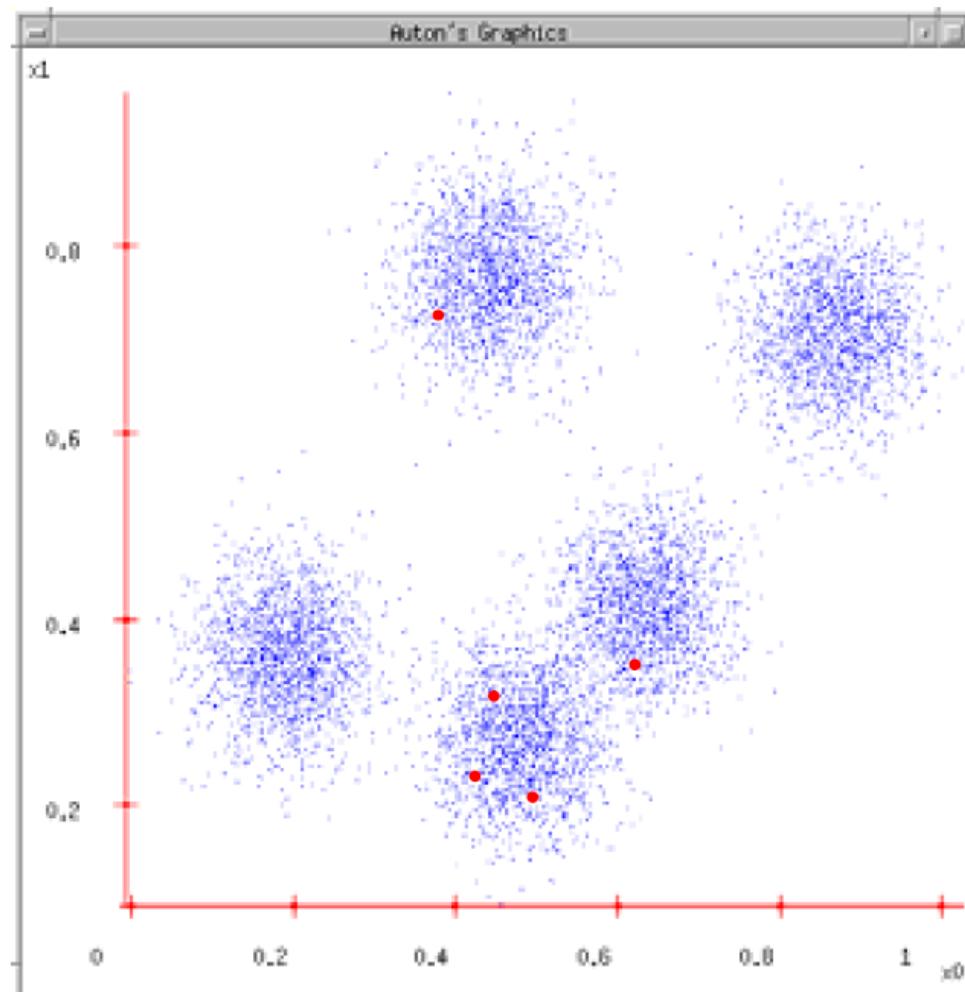
1. Ask user how many clusters they'd like.  
*(e.g.  $k=5$ )*



Source: A. Moore

# K-means

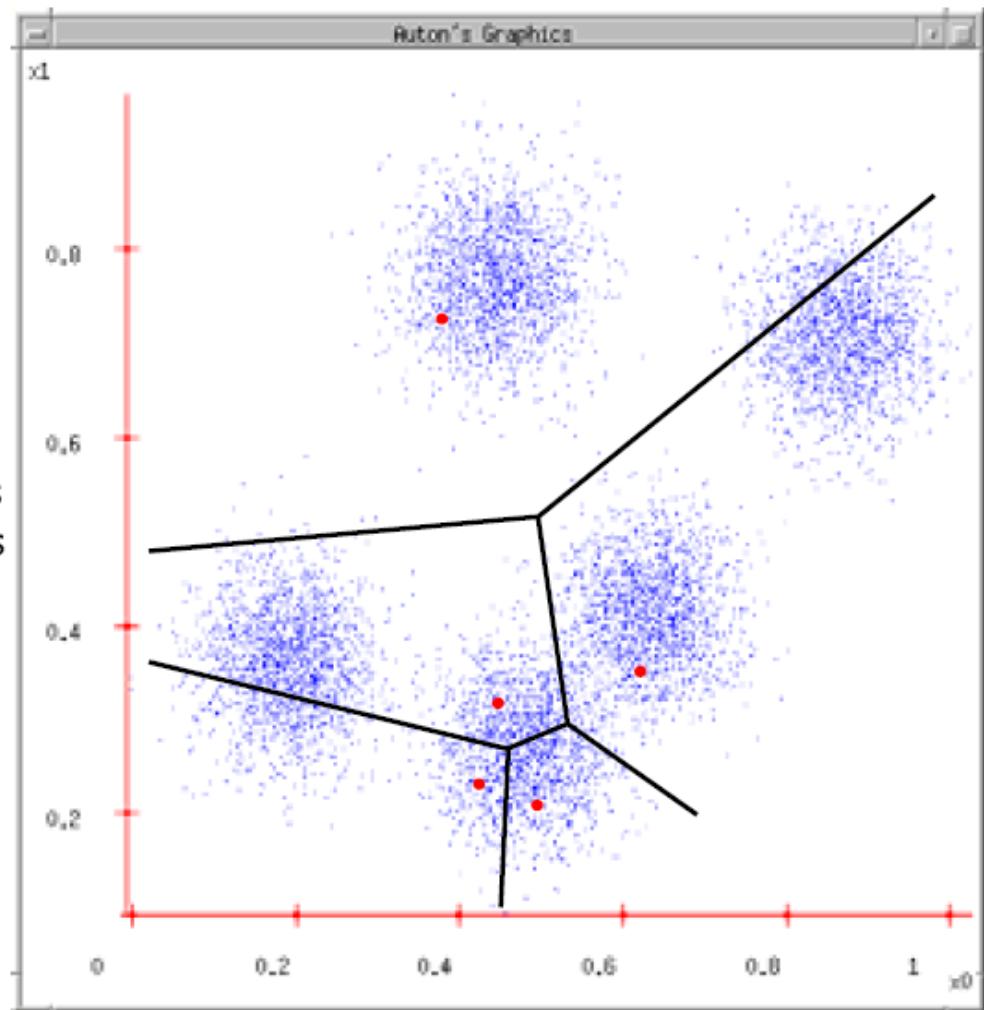
1. Ask user how many clusters they'd like.  
*(e.g. k=5)*
2. Randomly guess k cluster Center locations



Source: A. Moore

## K-means

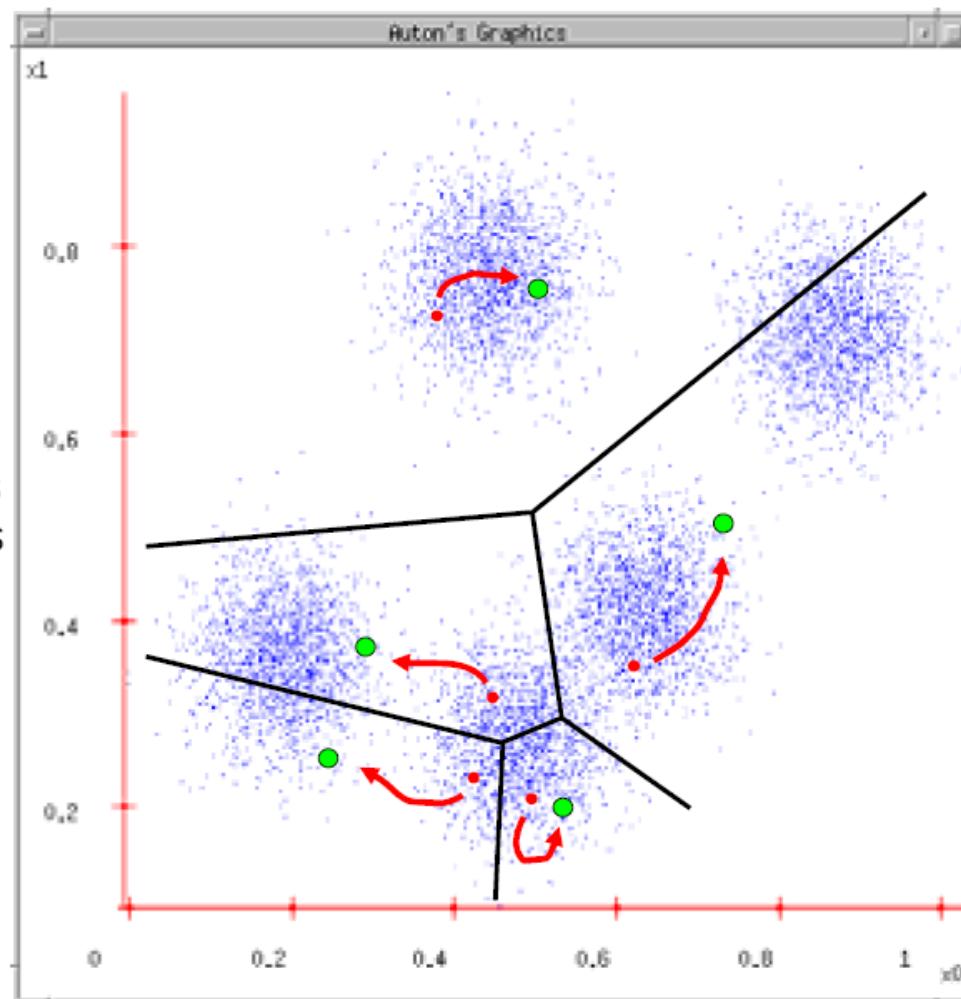
1. Ask user how many clusters they'd like.  
*(e.g. k=5)*
2. Randomly guess k cluster Center locations
3. Each datapoint finds out which Center it's closest to. (Thus each Center "owns" a set of datapoints)



Source: A. Moore

## K-means

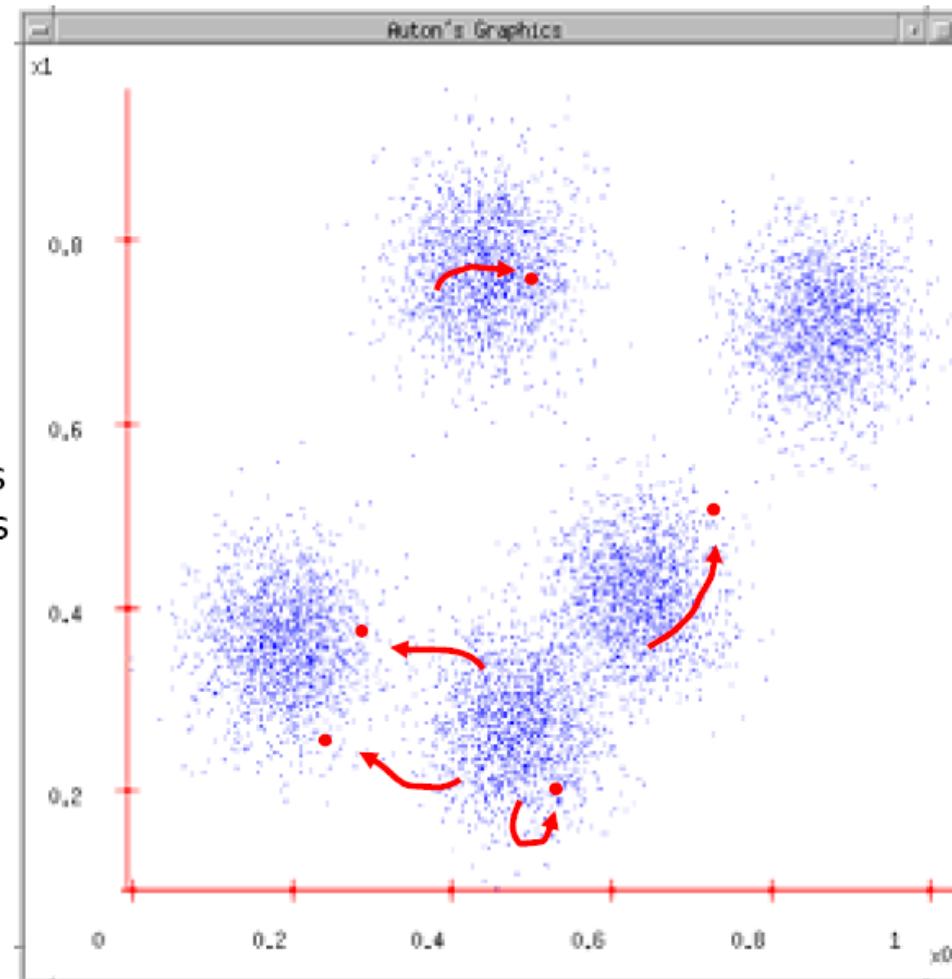
1. Ask user how many clusters they'd like.  
*(e.g. k=5)*
2. Randomly guess k cluster Center locations
3. Each datapoint finds out which Center it's closest to.
4. Each Center finds the centroid of the points it owns



Source: A. Moore

## K-means

1. Ask user how many clusters they'd like.  
*(e.g. k=5)*
2. Randomly guess k cluster Center locations
3. Each datapoint finds out which Center it's closest to.
4. Each Center finds the centroid of the points it owns...
5. ...and jumps there
6. ...Repeat until terminated!



Source: A. Moore

# Motivation: Interactive Clustering

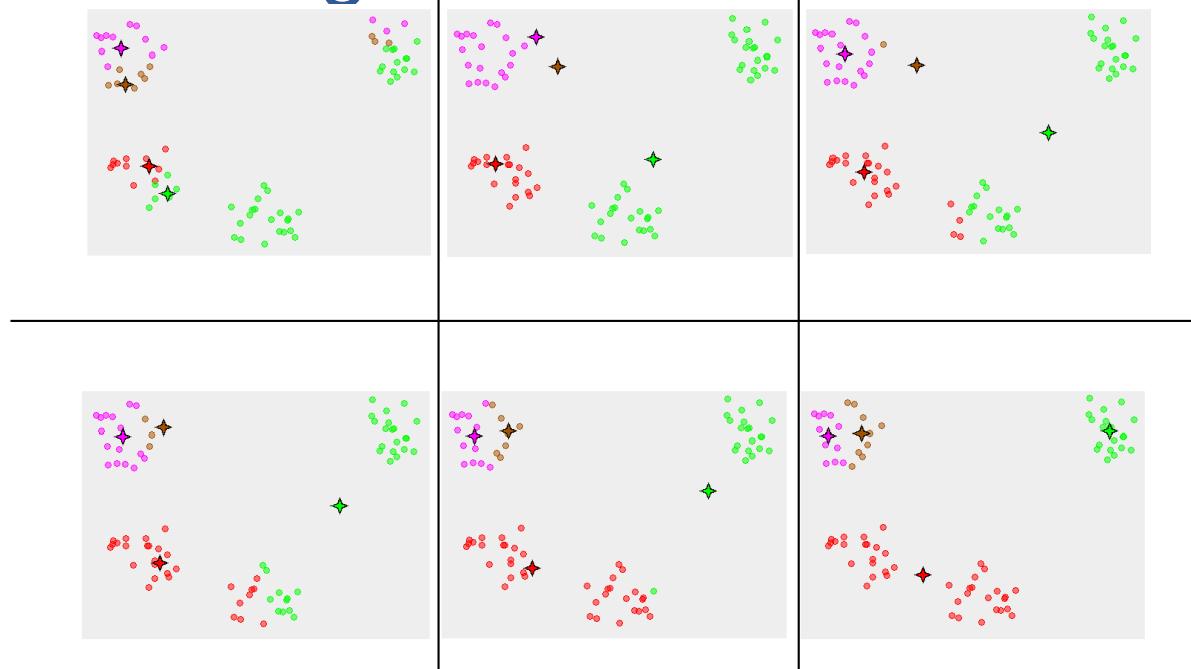


<https://www.naftaliharris.com/blog/visualizing-k-means-clustering/>



Explained Visually: <https://setosa.io/ev/>

## K-means converges to a local minimum



How can I try to fix this problem?

Adapted from James Hays

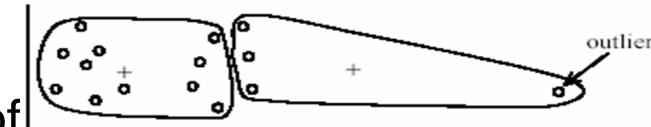
# K-means: pros and cons

## Pros

- Simple, fast to compute
- Converges to local minimum of within-cluster squared error

## Cons/issues

- Setting k?
  - One way: silhouette coefficient
- Sensitive to initial centers
  - Use heuristics or output of another method
- Sensitive to outliers
- Detects spherical clusters



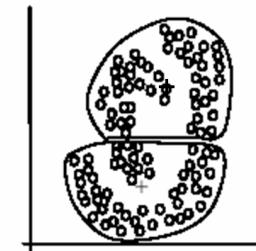
(A): Undesirable clusters



(B): Ideal clusters



(A): Two natural clusters

(B):  $k$ -means clusters

Adapted from K. Grauman

# Segmentation as Clustering

Depending on what we choose as the *feature space*, we can group pixels in different ways.

Grouping pixels based  
on **intensity** similarity



Feature space: intensity value (1-d)

Source: K. Grauman



K=2



K=3



Adapted from K. Grauman

# Segmentation as Clustering

Depending on what we choose as the *feature space*, we can group pixels in different ways.

Grouping pixels based  
on **intensity** similarity



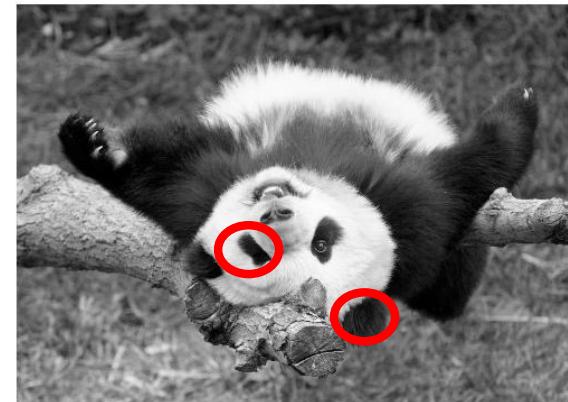
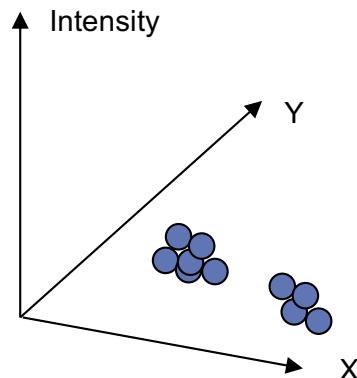
**Clusters** based on intensity  
similarity don't have to be  
spatially coherent.



# Segmentation as Clustering

Depending on what we choose as the *feature space*, we can group pixels in different ways.

Grouping pixels based  
on **intensity+position** similarity

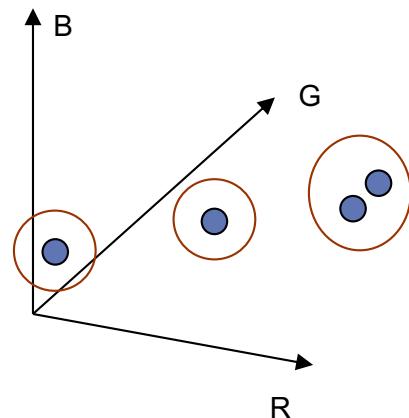


Both regions are black, but if we also include **position ( $x,y$ )**, then we could group the two into distinct segments; way to encode both similarity & proximity.

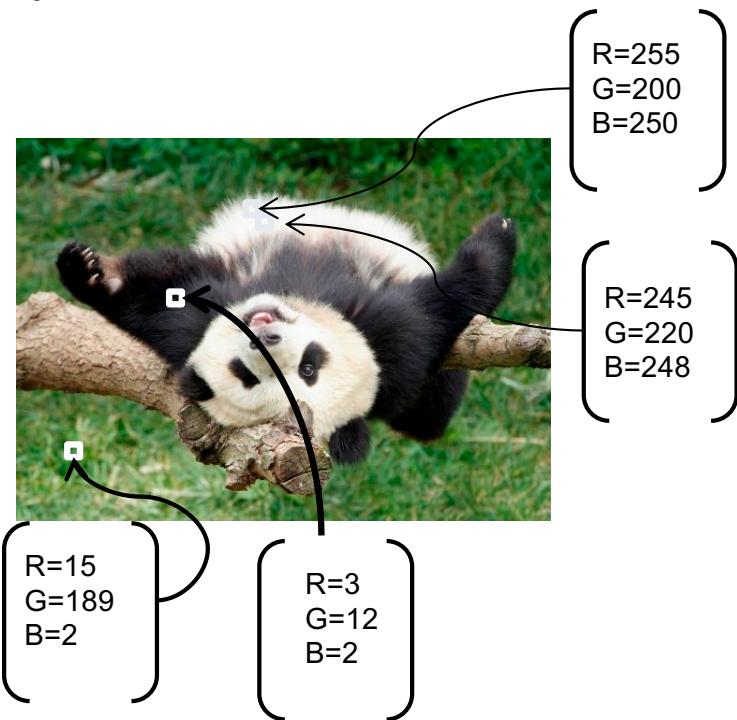
# Segmentation as Clustering

Depending on what we choose as the *feature space*, we can group pixels in different ways.

Grouping pixels based on **color** similarity



Feature space: color value (3-d)



Adapted from K. Grauman

# Segmentation as Clustering

- Color, brightness, position alone are not enough to distinguish all regions...

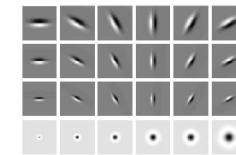
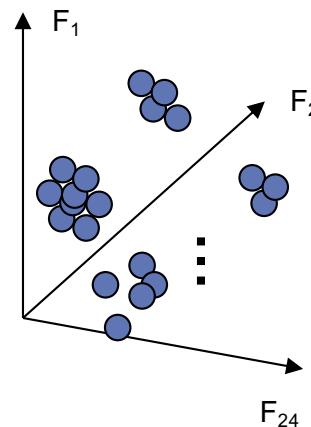


Source: L. Lazebnik

# Segmentation as Clustering

Depending on what we choose as the *feature space*, we can group pixels in different ways.

Grouping pixels based  
on **texture** similarity

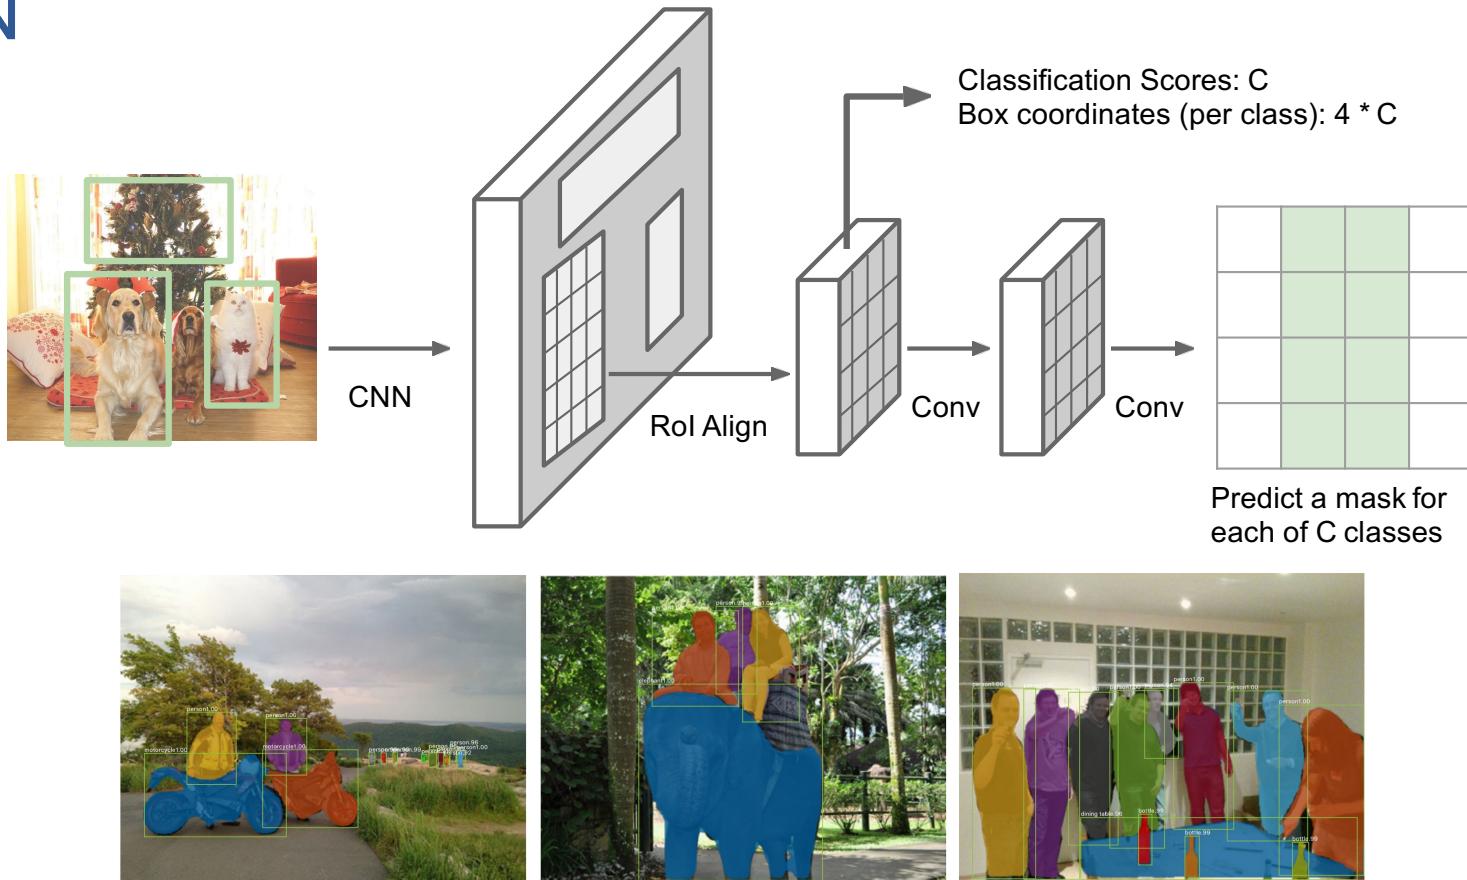


Filter bank  
of 24 filters

Feature space: filter bank responses (e.g., 24-d)

Source: K. Grauman

# State-of-the-art (instance) segmentation: Mask R-CNN



He et al, “[Mask R-CNN](#)”, ICCV 2017; slide adapted from Justin Johnson

# Lab 4: k-Means

Duration: 40 min

Use JPEG, PNG and GIF files less than 15 MB [[ahaslides](#)]



To join, go to: [ahaslides.com/58LUP](https://ahaslides.com/58LUP)



Please, run k-means on the provided dataset and upload your image result.

Get Feedback



0 0/100 ✓

# Summary

- **Segments:** use clustering (e.g. K-means) to group pixels by intensity, texture, etc.