

CS 1674: Visual Recognition

PhD. Nils Murrugarra-Llerena
nem177@pitt.edu



[Motivation] Visual Recognition

Medical Diagnostics: Aiding Pathologists

Imagine a pathologist examining a biopsy slide to detect cancer. This is a highly skilled, time-consuming, and mentally taxing task. They have to scan the entire slide under a microscope, looking for subtle cellular abnormalities that might indicate a malignant tumor. A single misdiagnosis can have devastating consequences.

Can we use CV to help pathologists in this task? How?



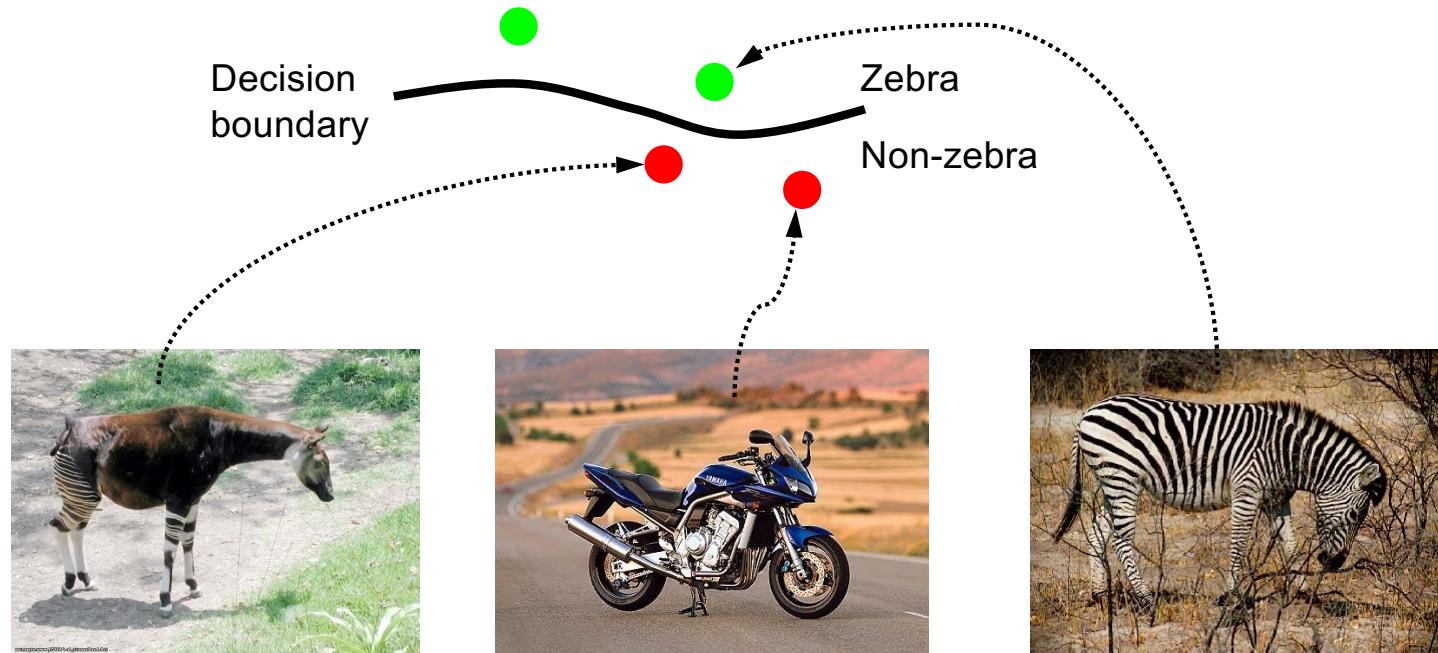
Plan for this lecture

- What is recognition?
 - a.k.a. classification, categorization
- Support vector machines
 - Separable case / non-separable case
 - Linear / non-linear (kernels)
- Example approach for scene classification
- Evaluation Metrics



Classification

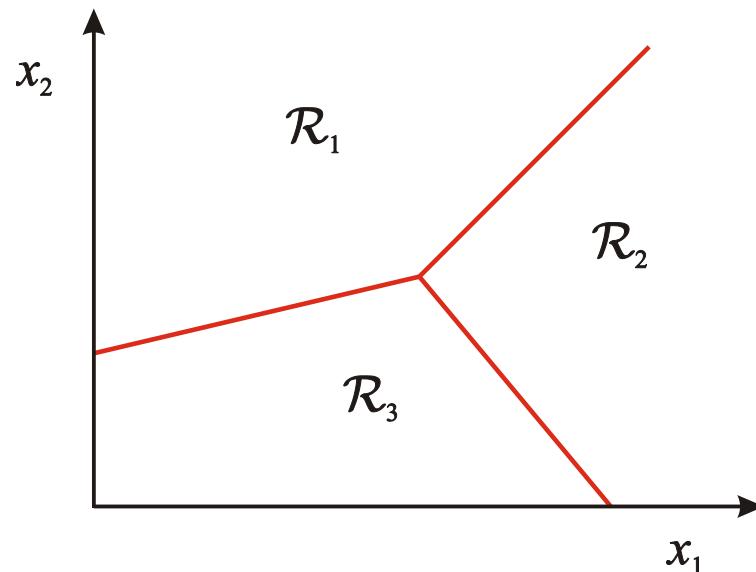
- Given a feature representation for images, how do we learn a model for distinguishing features from different classes?



Slide credit: L. Lazebnik

Classification

- Assign input vector to one of two or more classes
- Input space divided into *decision regions* separated by *decision boundaries*



Slide credit: L. Lazebnik

Examples of image classification

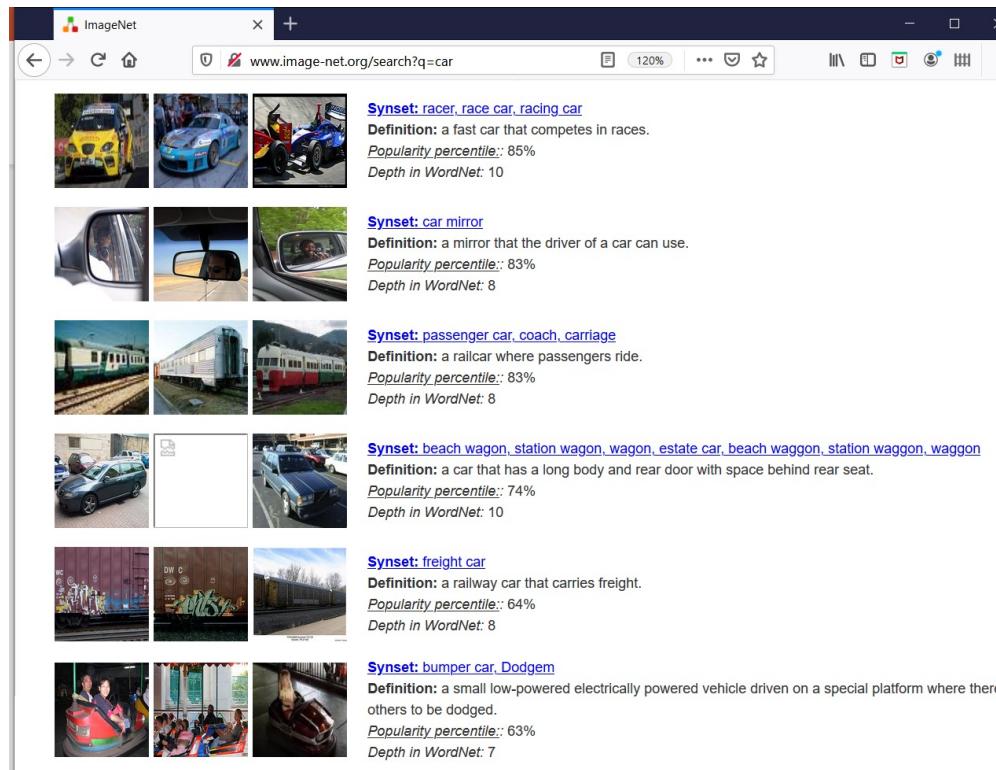
- Two-class (binary): Cat vs Dog



Adapted from D. Hoiem

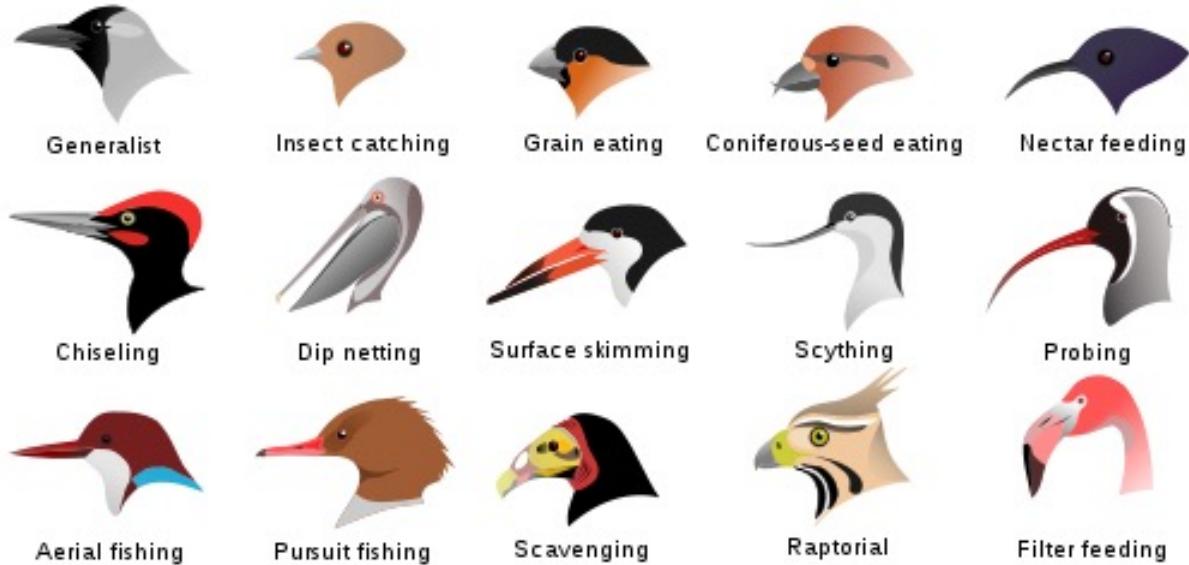
Examples of image classification

- Multi-class (often): Object recognition



Examples of image classification

- Fine-grained recognition



[Visipedia Project](#)

Slide credit: D. Hoiem

Examples of image classification

- Place recognition

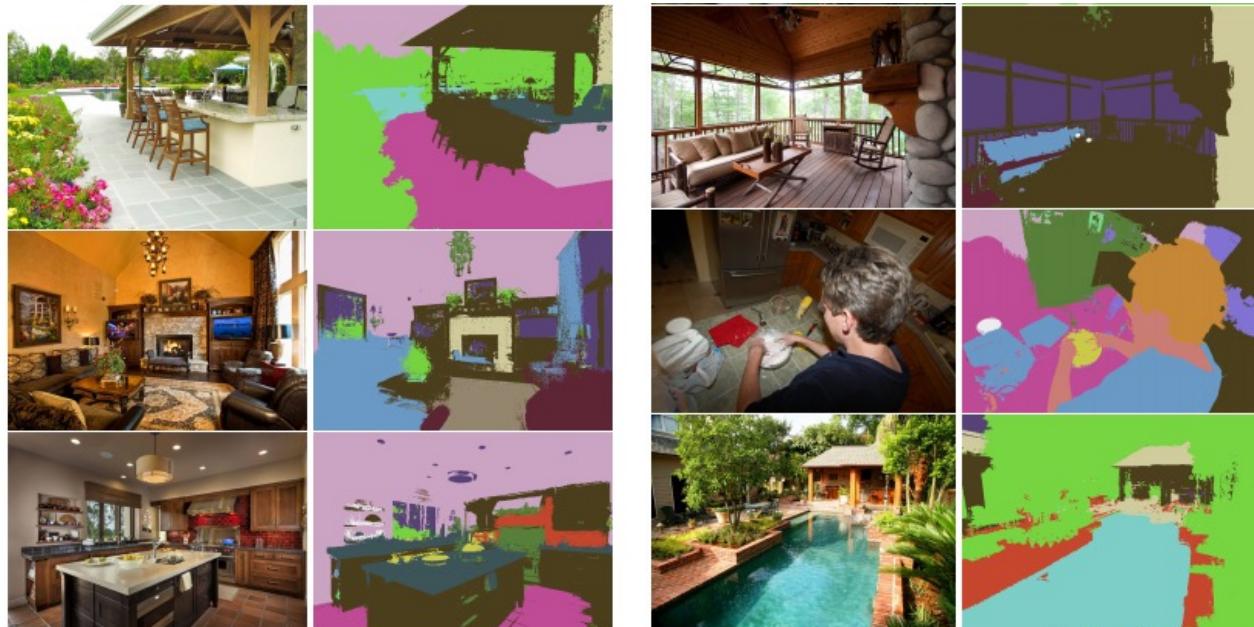


Places Database [[Zhou et al. NIPS 2014](#)]

Slide credit: D. Hoiem

Examples of image classification

- Material



[Bell et al. CVPR 2015]

Slide credit: D. Hoiem

Examples of image classification

- Dating historical photos



1940

1953

1966

1977

[[Palermo et al. ECCV 2012](#)]

Slide credit: D. Hoiem

Examples of image classification

- Image style recognition



HDR



Macro



Baroque



Rococo



Vintage



Noir



Northern Renaissance



Cubism



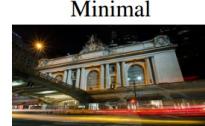
Minimal



Hazy



Impressionism



Long Exposure



Romantic



Abs. Expressionism



Color Field Painting

Flickr Style: 80K images covering 20 styles.

Wikipaintings: 85K images for 25 art genres.

[[Karayev et al. BMVC 2014](#)]

Slide credit: D. Hoiem

Recognition: An Image Classifier

```
def classify_image(image):  
    # Some magic here?  
    return class_label
```



Unlike e.g. sorting a list of numbers,

no obvious way to hard-code the algorithm for recognizing
a cat, or other classes.

Recognition: A machine learning approach



Recognition: A machine learning approach

1. Collect a dataset of images and labels
2. Use Machine Learning algorithms to train a classifier
3. Evaluate the classifier on new images

```
def train(images, labels):
    # Machine learning!
    return model
```

```
def predict(model, test_images):
    # Use model to predict labels
    return test_labels
```

Example training set

airplane



automobile



bird

cat

deer

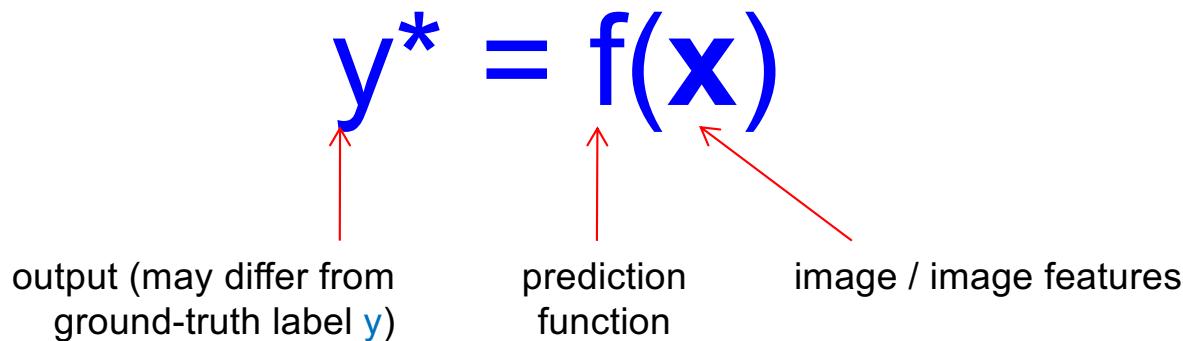
Slide Credit: <https://cs231n.stanford.edu/>

The machine learning framework

- Apply a prediction function to a feature representation of the image to get the desired output:

$$f(\text{apple}) = \text{"apple"}$$
$$f(\text{tomato}) = \text{"tomato"}$$
$$f(\text{cow}) = \text{"cow"}$$

The machine learning framework

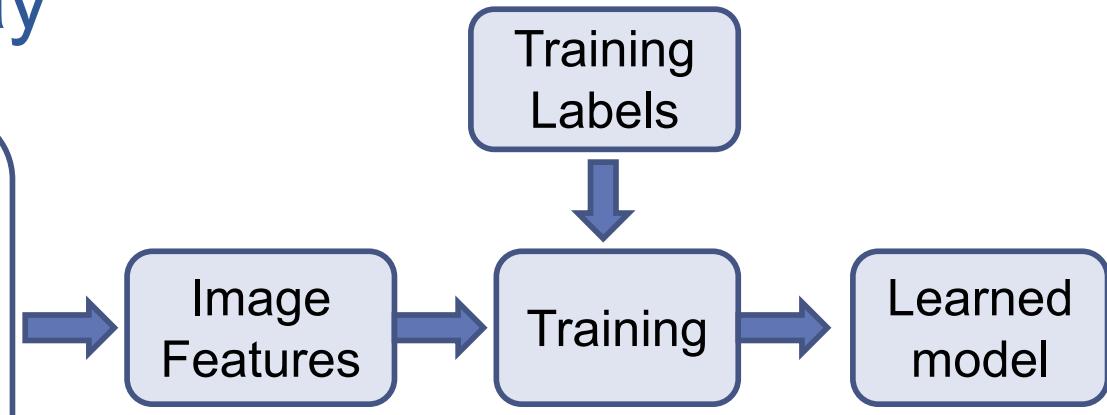
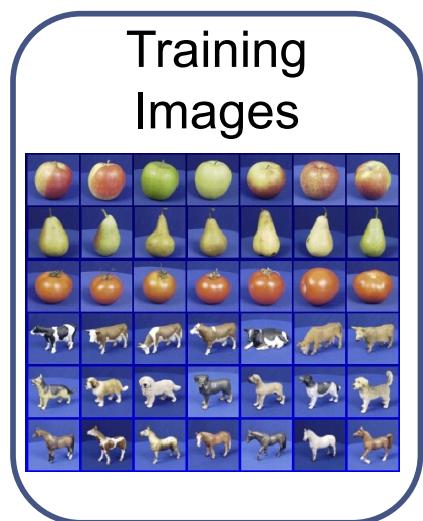
$$y^* = f(x)$$


output (may differ from ground-truth label y) prediction function image / image features

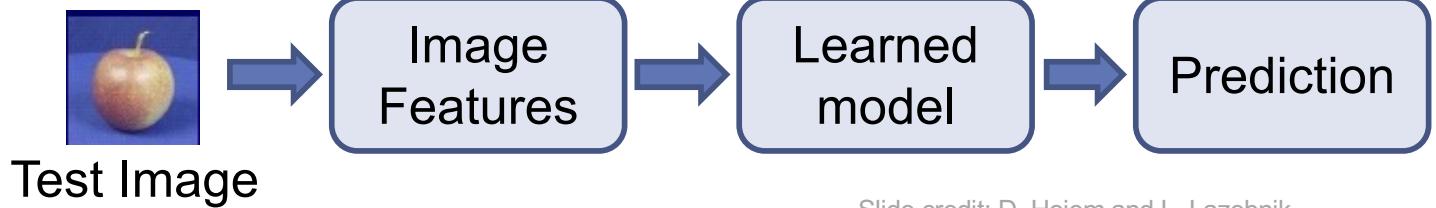
- **Training:** given a *training set* of labeled examples $\{(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_N, y_N)\}$, estimate the prediction function f by minimizing the prediction error on the training set, e.g. $|f(\mathbf{x}_i) - y_i|$
 - Evaluate multiple hypotheses $f_1, f_2, f_H \dots$ and pick the best one as f
- **Testing:** apply f to a never-before-seen *test example* \mathbf{x} and output the predicted value $y^* = f(\mathbf{x})$

The old-school way

Training

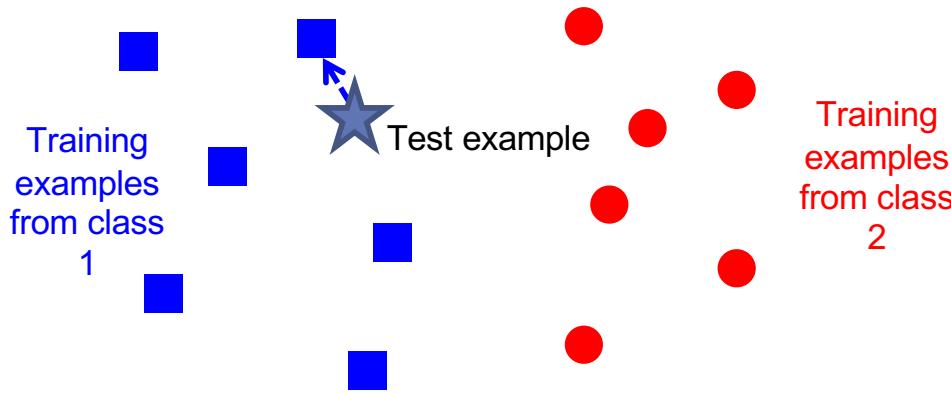


Testing



Slide credit: D. Hoiem and L. Lazebnik

The simplest classifier: Nearest Neighbor Classifier

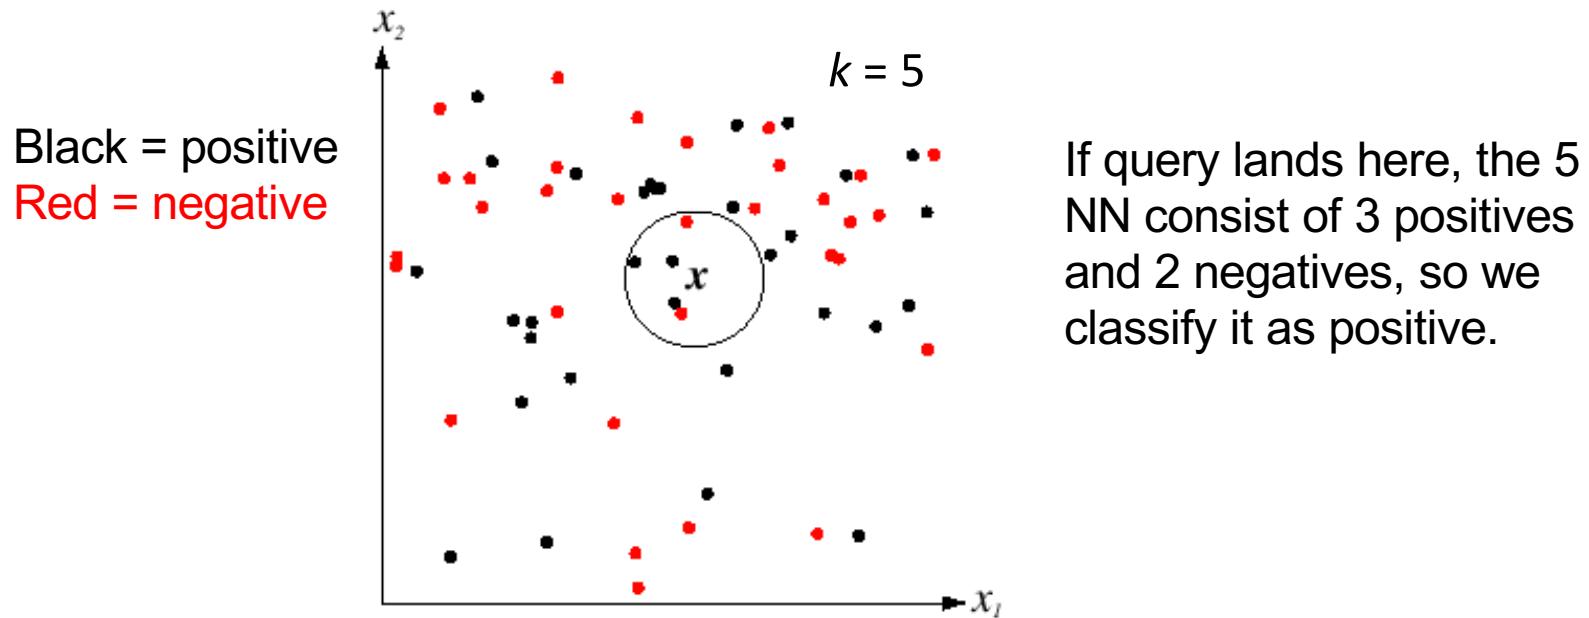


$f(\mathbf{x})$ = label of the training example nearest to \mathbf{x}

- All we need is a distance function for our inputs
- No training required!

K-Nearest Neighbors classification

- For a new point, find the k closest points from training data
- Labels of the k points “vote” to classify



Slide credit: D. Lowe

K-Nearest Neighbor - Summary

```
def train(images, labels):  
    # Machine learning!  
    return model
```

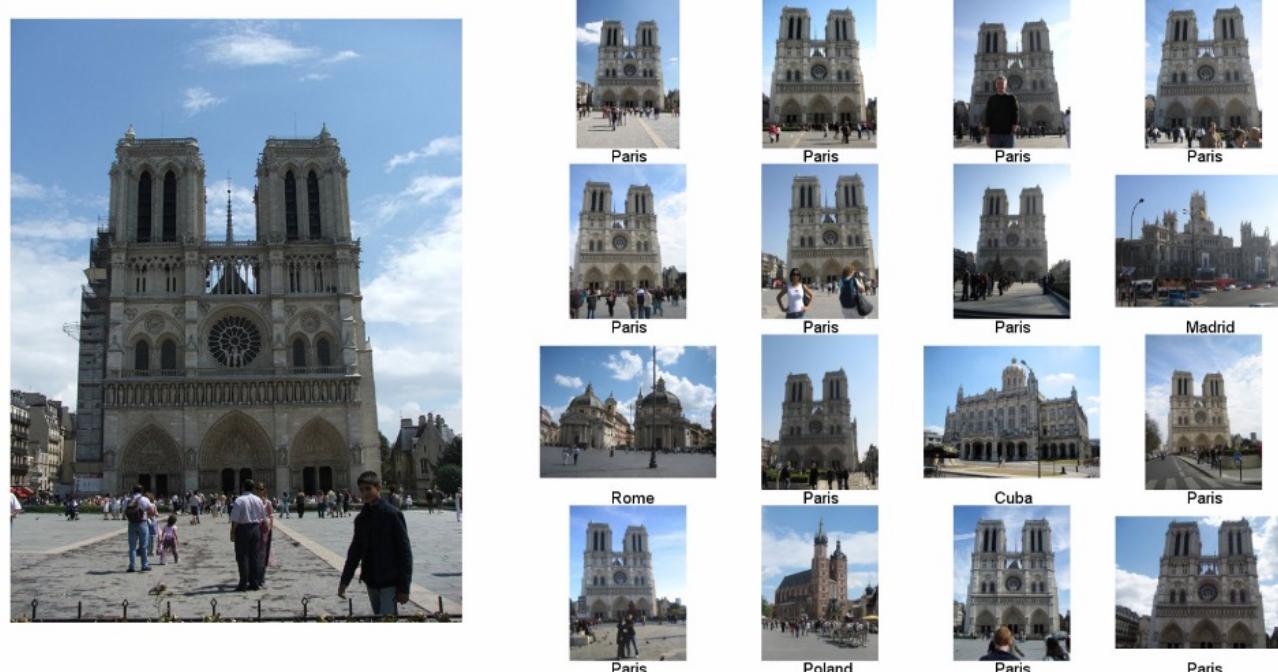
Memorize all data and labels

```
def predict(model, test_images):  
    # Use model to predict labels  
    return test_labels
```

Predict the label of the
most similar training
image

[Application]Im2gps: Estimating Geographic Information from a Single Image [James Hays and Alexei Efros, CVPR 2008]

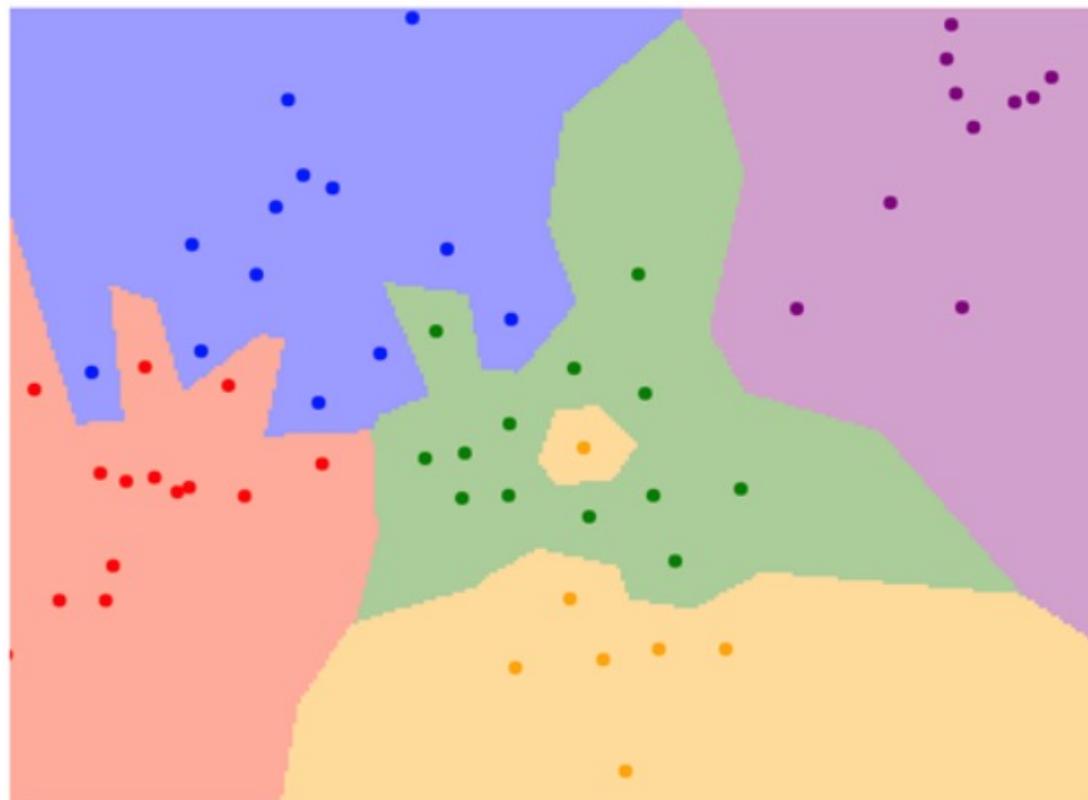
Where was this image taken?



Nearest Neighbors according to BOW-SIFT + color histogram + a few others

Slide credit: James Hays

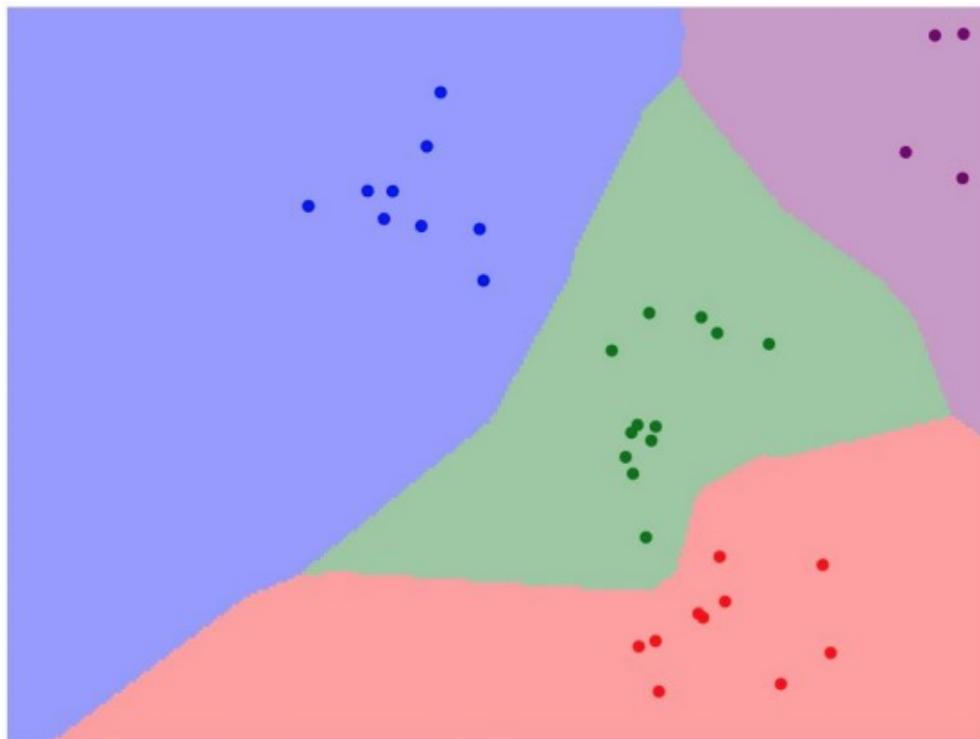
K-Nearest Neighbor - Visualization



1-nearest
Neighbor

Slide Credit: <https://cs231n.stanford.edu/>

K-Nearest Neighbor – Interact – Try it yourself



<http://vision.stanford.edu/teaching/cs231n-demos/knn/>

Slide Credit: <https://cs231n.stanford.edu/>

Setting Hyperparameters: Best value of k?

Idea #1: Choose hyperparameters that work best on the training data

train

Setting Hyperparameters: Best value of k?

train

Idea #4: Cross-Validation: Split data into folds, try each fold as validation and average the results

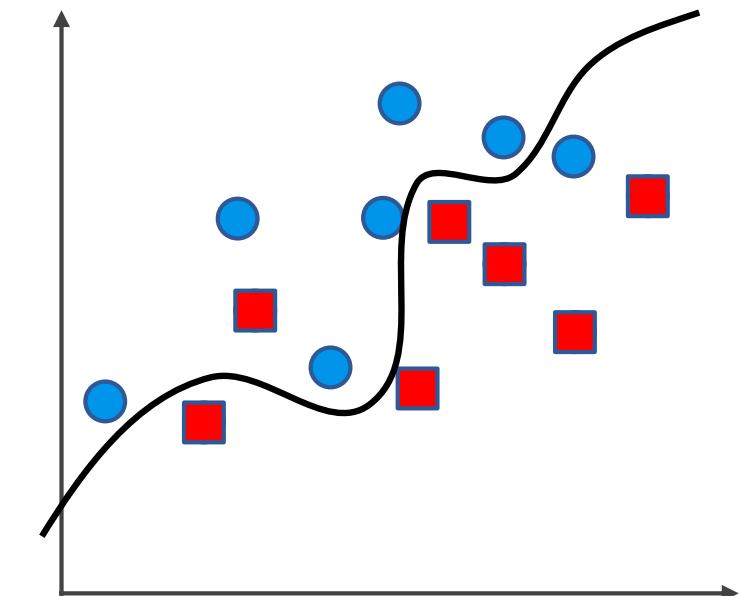
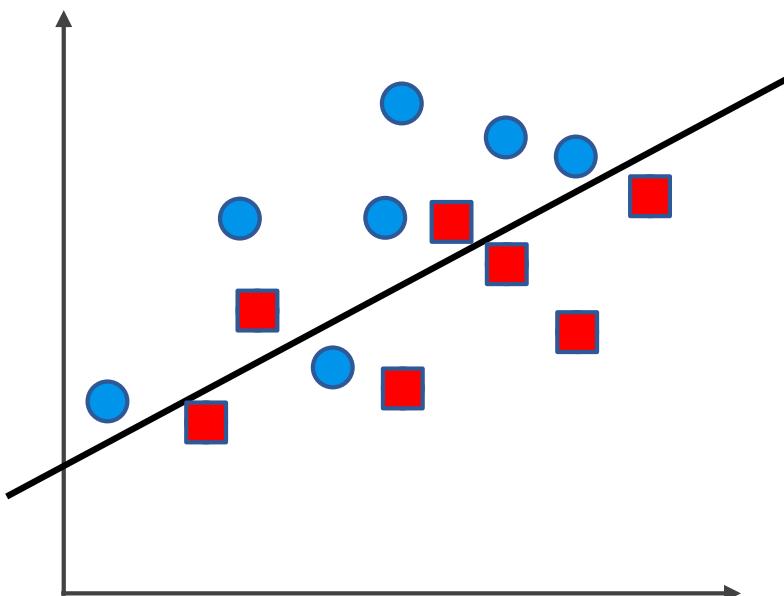
fold 1	fold 2	fold 3	fold 4	fold 5
fold 1	fold 2	fold 3	fold 4	fold 5
fold 1	fold 2	fold 3	fold 4	fold 5
fold 1	fold 2	fold 3	fold 4	fold 5
fold 1	fold 2	fold 3	fold 4	fold 5

test

Useful for small datasets, but not used too frequently in deep learning

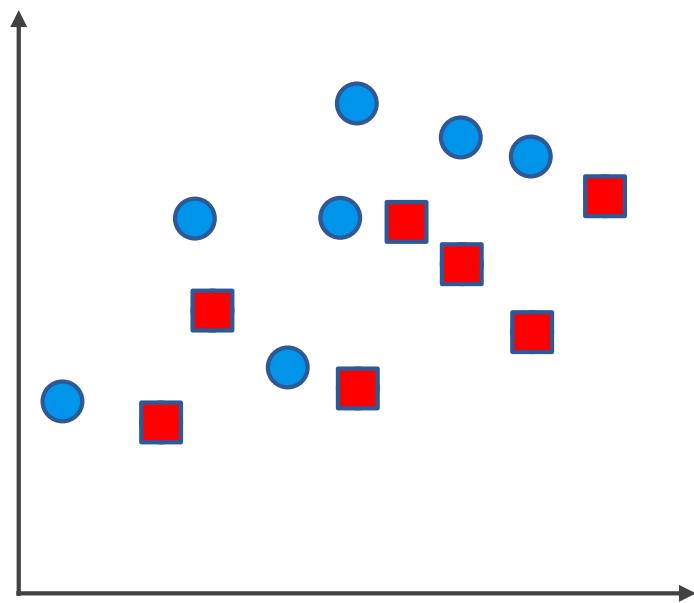
Slide Credit: <https://cs231n.stanford.edu/>

Which model is better? Why validating?



Slide Credit: Prof. Sandra Avila - UNICAMP

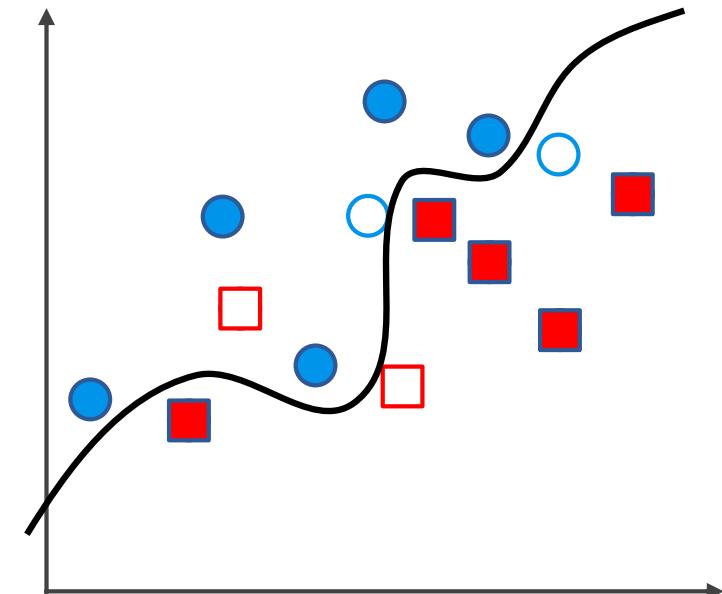
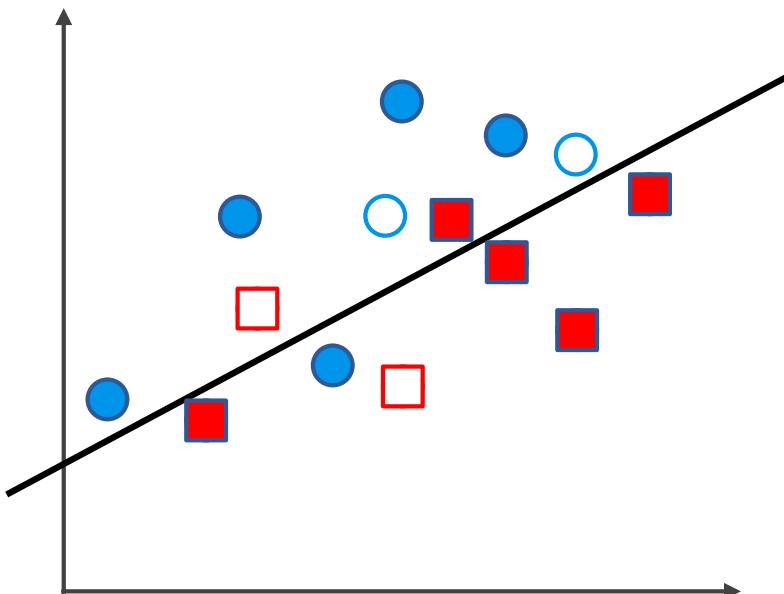
Why validating?



Slide Credit: Prof. Sandra Avila - UNICAMP



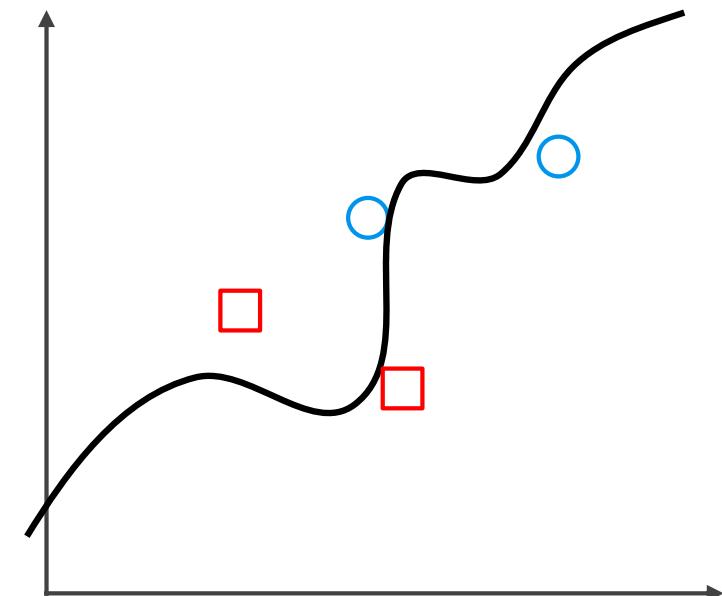
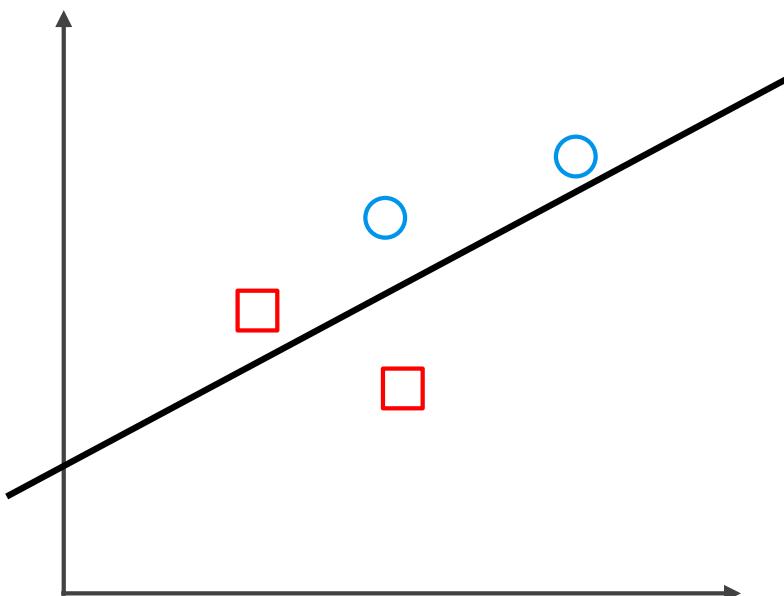
Why validating?



Slide Credit: Prof. Sandra Avila - UNICAMP

- ■ Training
- □ Validation

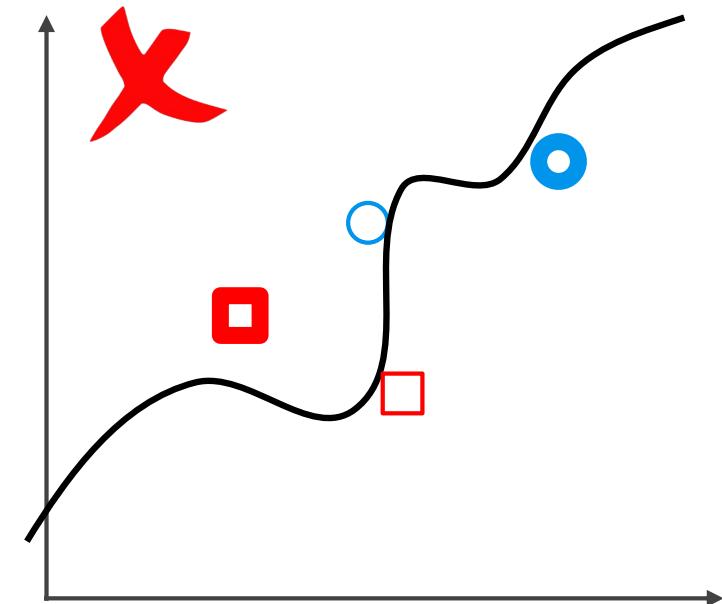
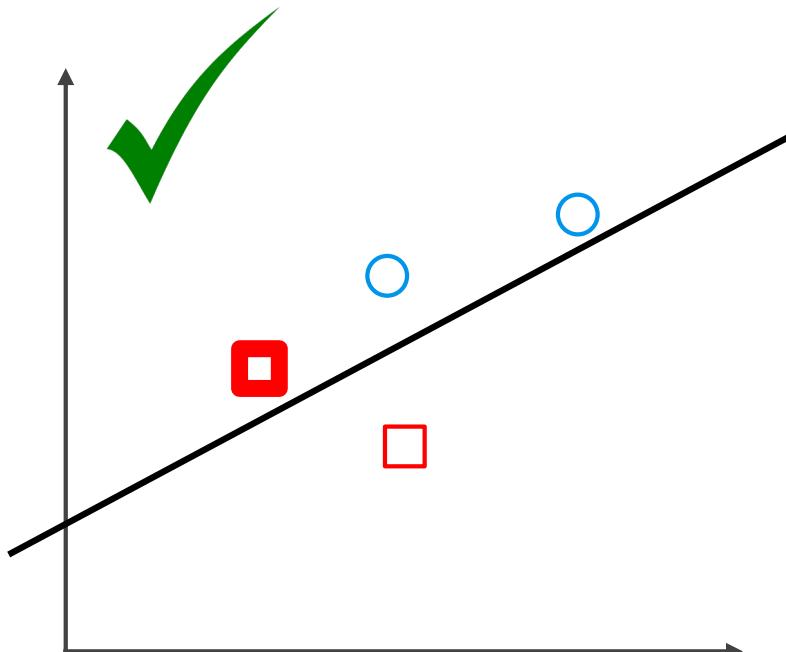
Why validating?



Slide Credit: Prof. Sandra Avila - UNICAMP

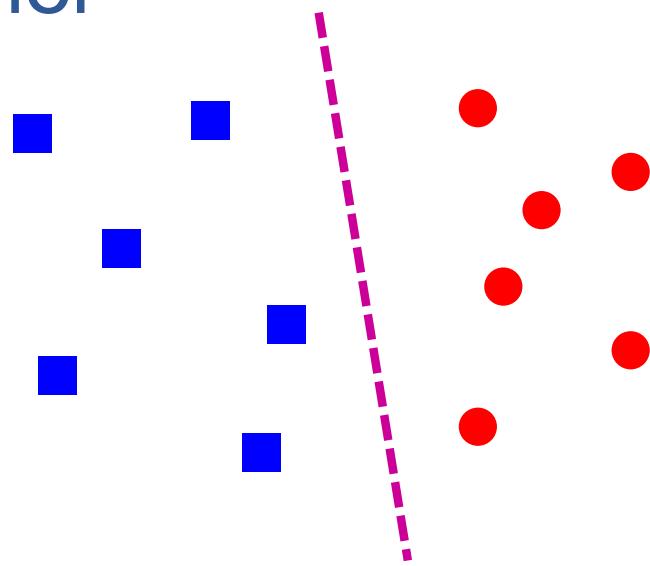


Why validating?



Slide Credit: Prof. Sandra Avila - UNICAMP

Linear classifier

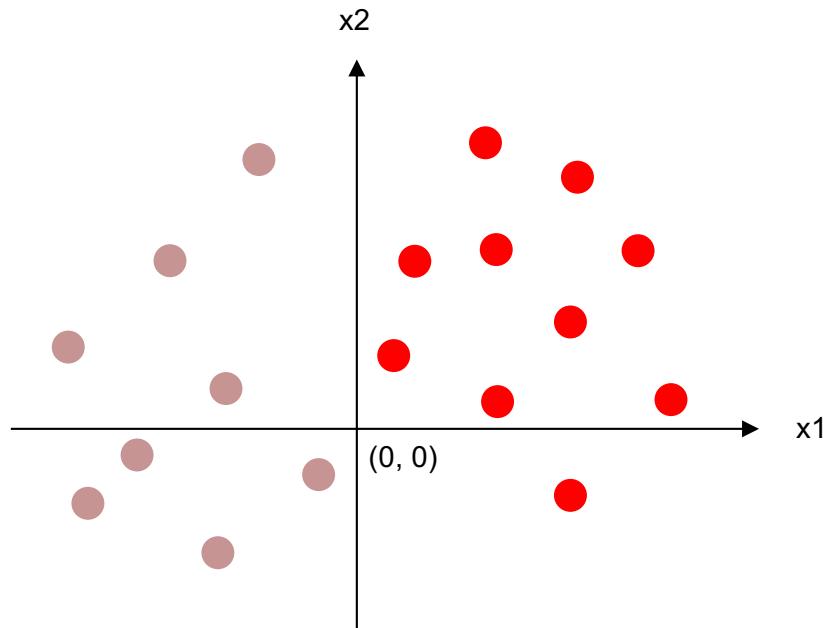


- Find a *linear function* to separate the classes

$$f(\mathbf{x}) = \text{sgn}(w_1x_1 + w_2x_2 + \dots + w_Dx_D) = \text{sgn}(\mathbf{w} \cdot \mathbf{x})$$

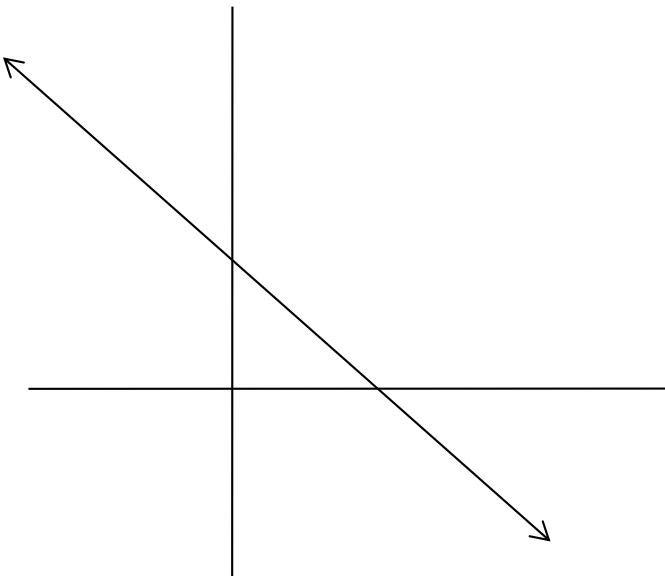
Linear Classifier

- Decision = $\text{sign}(\mathbf{w}^T \mathbf{x}) = \text{sign}(w_1 * x_1 + w_2 * x_2)$



- What should the weights be?

Lines in \mathbb{R}^2



Let $\mathbf{w} = \begin{bmatrix} a \\ c \end{bmatrix}$ $\mathbf{x} = \begin{bmatrix} x \\ y \end{bmatrix}$

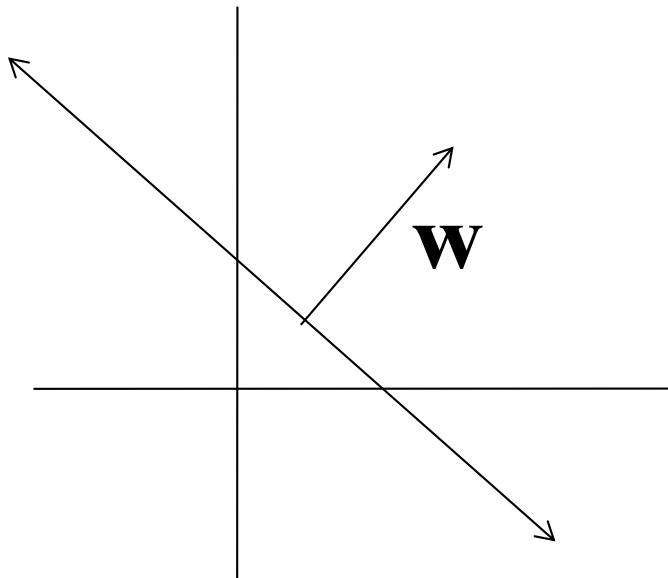
$$ax + cy + b = 0$$

Compare to:
*slope**x + *y-intercept* = y

$$\begin{aligned} ax + b &= -cy \\ (-a/c)x + (-b/c) &= y \end{aligned}$$

Slope: $-a/c$
Y-intercept: $-b/c$

Lines in \mathbb{R}^2



Slope: $-a/c$
Y-intercept: $-b/c$

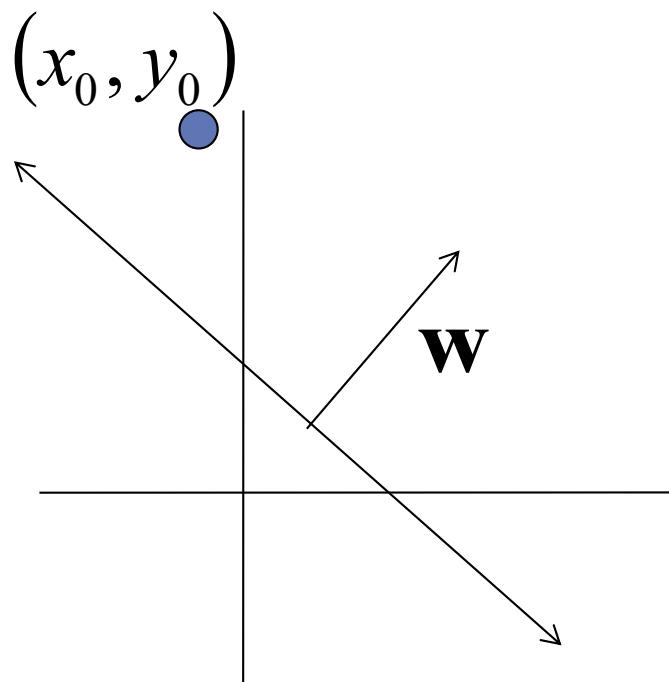
Let $\mathbf{w} = \begin{bmatrix} a \\ c \end{bmatrix}$ $\mathbf{x} = \begin{bmatrix} x \\ y \end{bmatrix}$

$$ax + cy + b = 0$$



$$\mathbf{w} \cdot \mathbf{x} + b = 0$$

Lines in \mathbb{R}^2



Slope: $-a/c$
 Y-intercept: $-b/c$

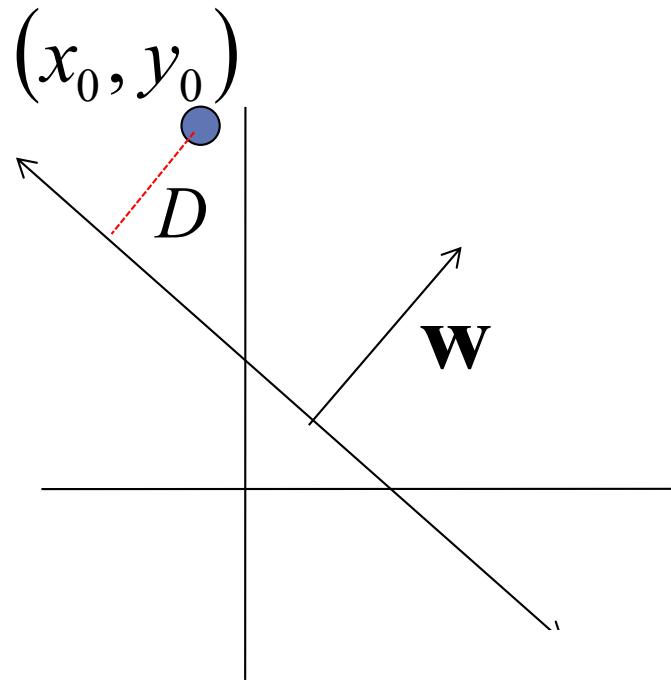
Let $\mathbf{w} = \begin{bmatrix} a \\ c \end{bmatrix}$ $\mathbf{x} = \begin{bmatrix} x \\ y \end{bmatrix}$

$$ax + cy + b = 0$$

↔

$$\mathbf{w} \cdot \mathbf{x} + b = 0$$

Lines in \mathbb{R}^2



$$D = \frac{|ax_0 + cy_0 + b|}{\sqrt{a^2 + c^2}}$$

Slope: $-a/c$
Y-intercept: $-b/c$

Let $\mathbf{w} = \begin{bmatrix} a \\ c \end{bmatrix}$ $\mathbf{x} = \begin{bmatrix} x \\ y \end{bmatrix}$

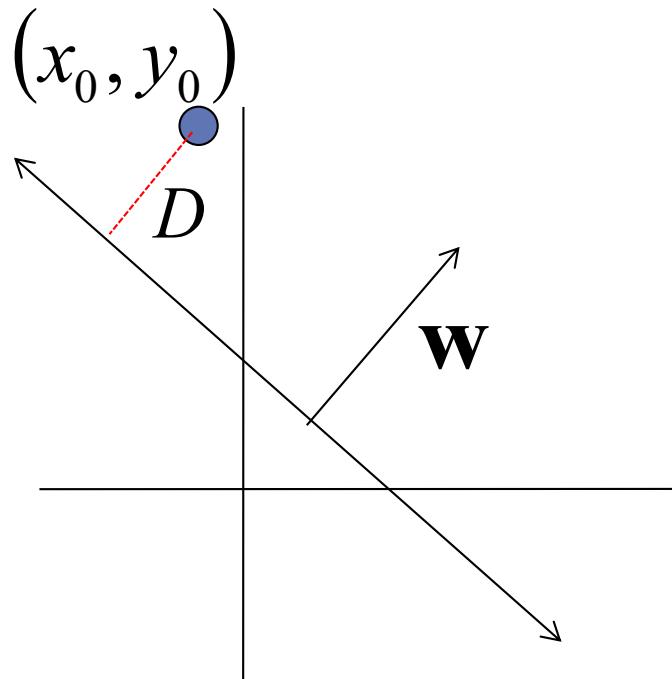
$$ax + cy + b = 0$$



$$\mathbf{w} \cdot \mathbf{x} + b = 0$$

} distance from
point to line

Lines in \mathbb{R}^2



$$D = \frac{|ax_0 + cy_0 + b|}{\sqrt{a^2 + c^2}} = \frac{|\mathbf{w}^\top \mathbf{x} + b|}{\|\mathbf{w}\|}$$

Slope: $-a/c$
Y-intercept: $-b/c$

Let $\mathbf{w} = \begin{bmatrix} a \\ c \end{bmatrix}$ $\mathbf{x} = \begin{bmatrix} x \\ y \end{bmatrix}$

$$ax + cy + b = 0$$

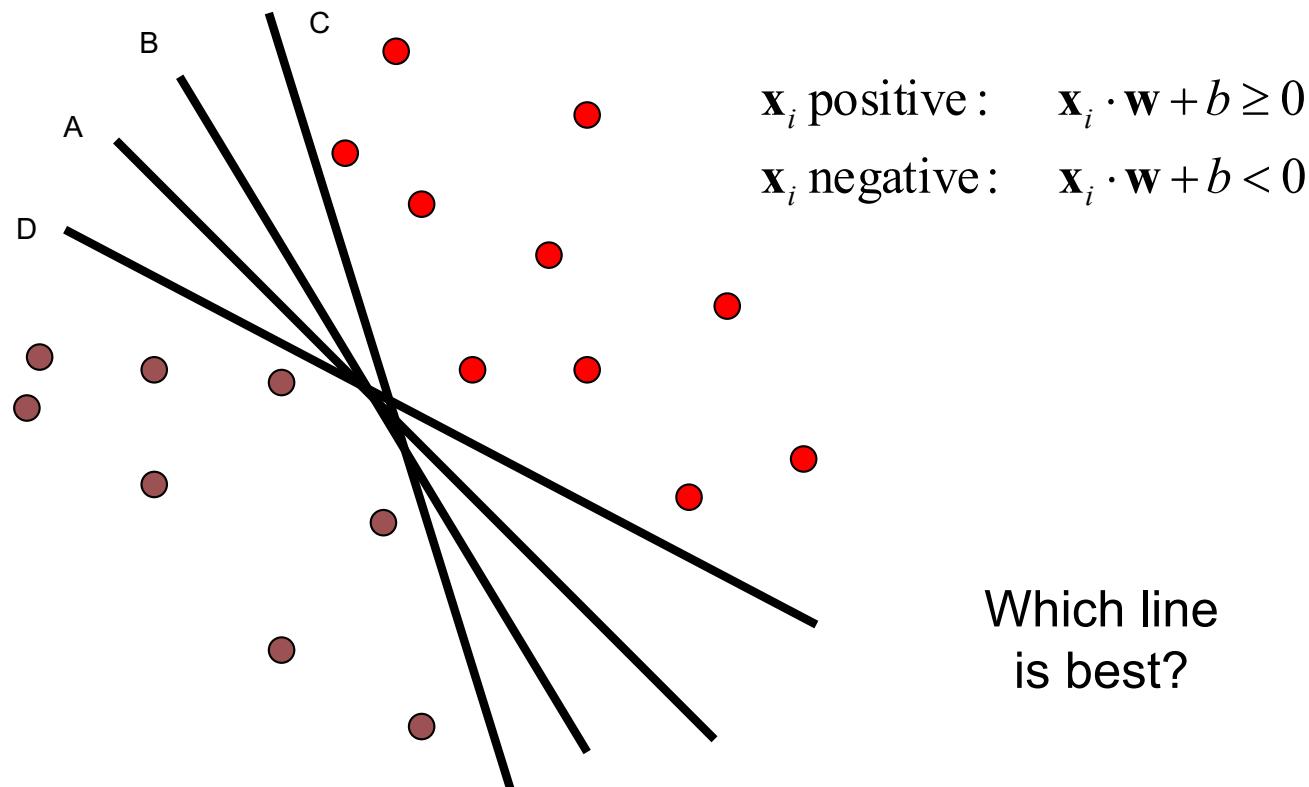


$$\mathbf{w} \cdot \mathbf{x} + b = 0$$

distance from
point to line

Linear classifiers

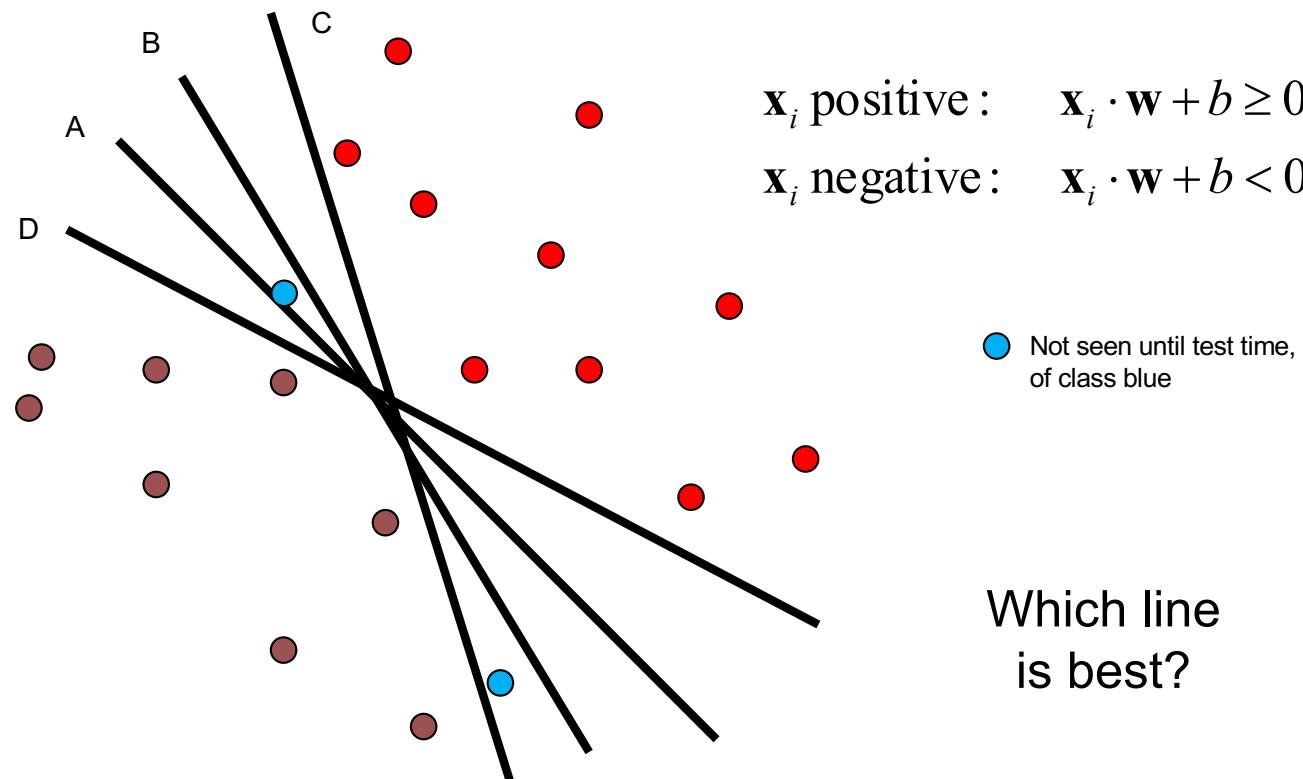
- Find linear function to separate positive and negative examples



C. Burges, [A Tutorial on Support Vector Machines for Pattern Recognition](#), Data Mining and Knowledge Discovery, 1998

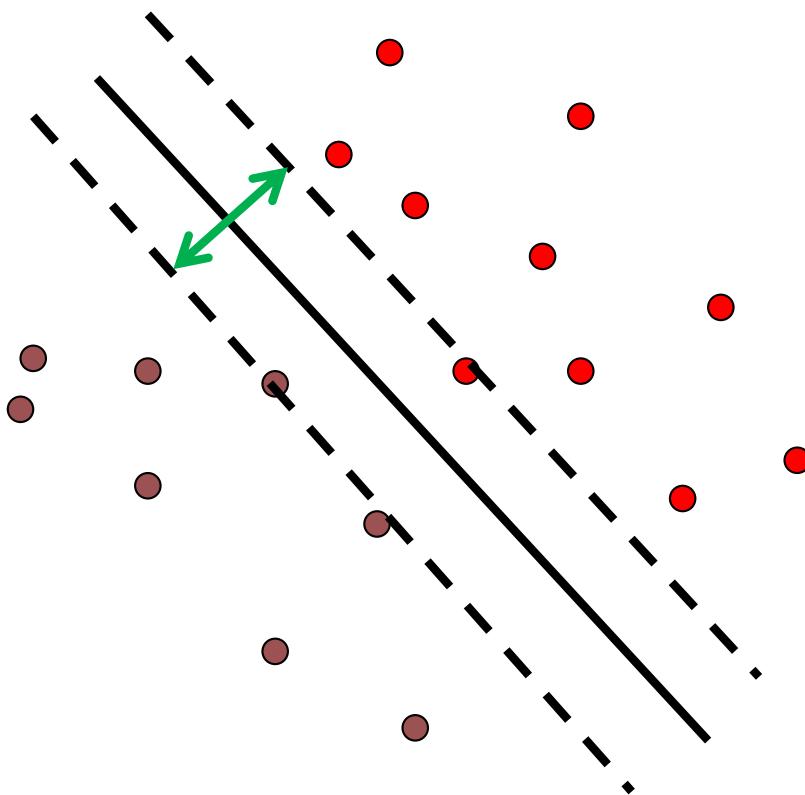
Linear classifiers

- Find linear function to separate positive and negative examples



C. Burges, [A Tutorial on Support Vector Machines for Pattern Recognition](#), Data Mining and Knowledge Discovery, 1998

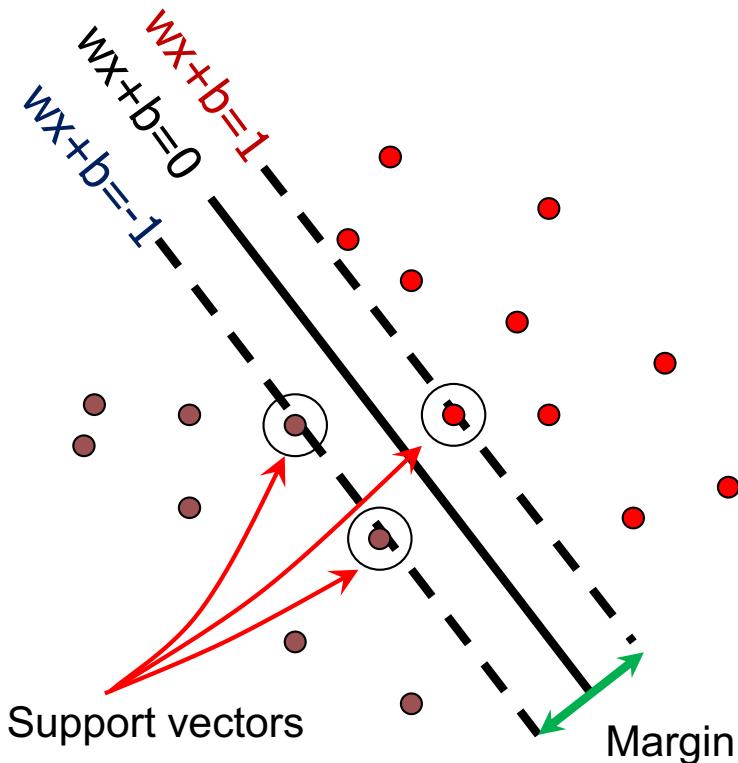
Support vector machines



- Discriminative classifier based on *optimal separating line* (for 2d case)
- Maximize the *margin* between the positive and negative training examples

Support vector machines

- Want line that maximizes the margin.



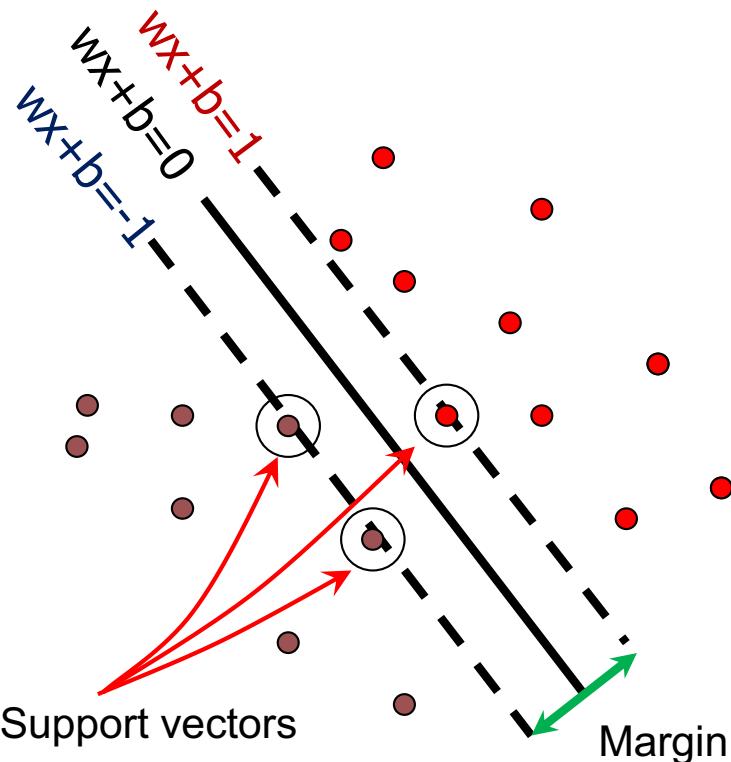
\mathbf{x}_i positive ($y_i = 1$): $\mathbf{x}_i \cdot \mathbf{w} + b \geq 1$

\mathbf{x}_i negative ($y_i = -1$): $\mathbf{x}_i \cdot \mathbf{w} + b \leq -1$

For support, vectors, $\mathbf{x}_i \cdot \mathbf{w} + b = \pm 1$

Support vector machines

- Want line that maximizes the margin.



\mathbf{x}_i positive ($y_i = 1$): $\mathbf{x}_i \cdot \mathbf{w} + b \geq 1$

\mathbf{x}_i negative ($y_i = -1$): $\mathbf{x}_i \cdot \mathbf{w} + b \leq -1$

For support, vectors, $\mathbf{x}_i \cdot \mathbf{w} + b = \pm 1$

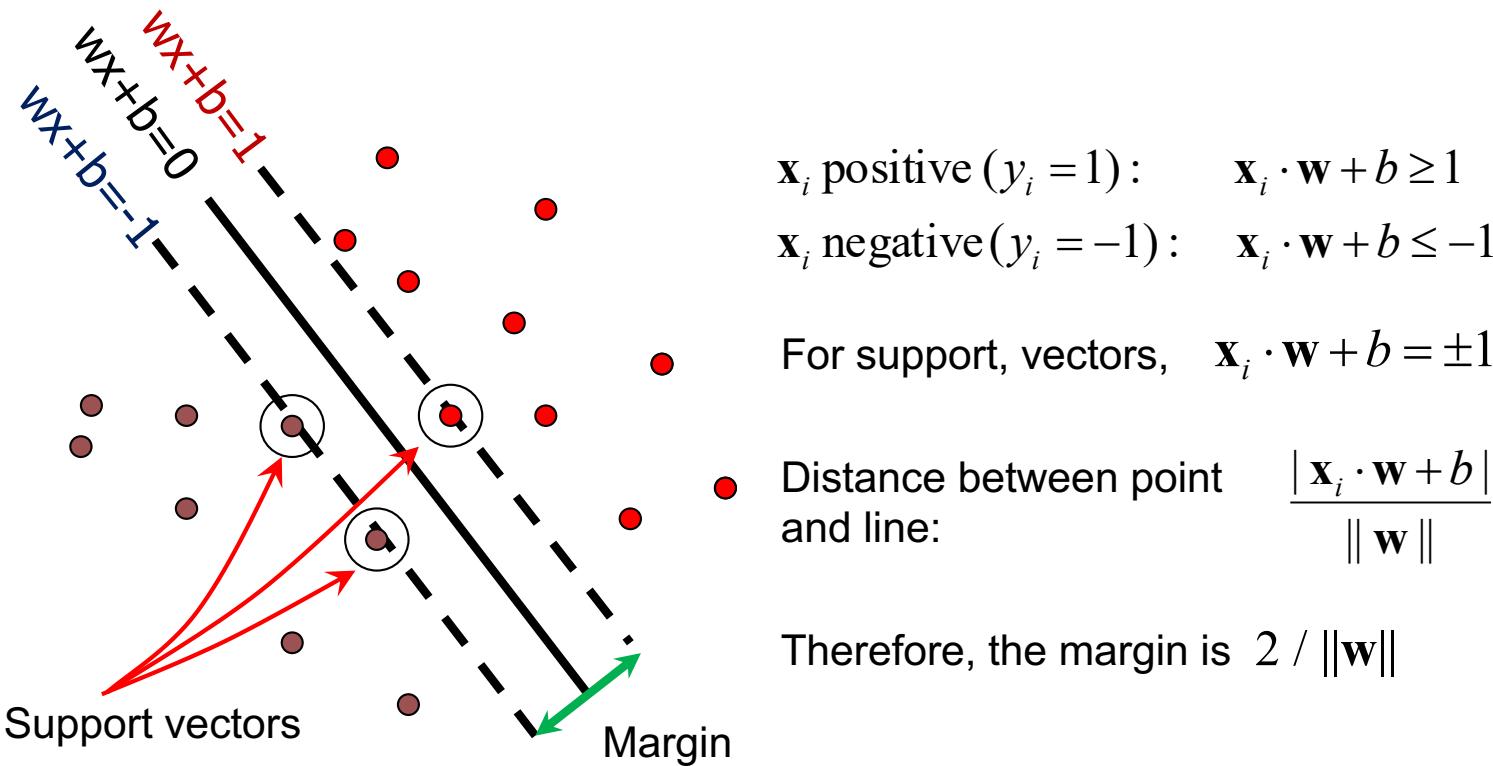
Distance between point and line:
$$\frac{|\mathbf{x}_i \cdot \mathbf{w} + b|}{\|\mathbf{w}\|}$$

For support vectors:

$$\frac{\mathbf{w}^T \mathbf{x}_i + b}{\|\mathbf{w}\|} = \frac{\pm 1}{\|\mathbf{w}\|} \quad M = \left| \frac{1}{\|\mathbf{w}\|} - \frac{-1}{\|\mathbf{w}\|} \right| = \frac{2}{\|\mathbf{w}\|}$$

Support vector machines

- Want line that maximizes the margin.



Finding the maximum margin line

1. Maximize margin $2/\|\mathbf{w}\|$
2. Correctly classify all training data points:

$$\mathbf{x}_i \text{ positive } (y_i = 1) : \quad \mathbf{x}_i \cdot \mathbf{w} + b \geq 1$$

$$\mathbf{x}_i \text{ negative } (y_i = -1) : \quad \mathbf{x}_i \cdot \mathbf{w} + b \leq -1$$

- *Quadratic optimization problem:*

-

Minimize $\frac{1}{2} \mathbf{w}^T \mathbf{w}$

Subject to $y_i(\mathbf{w} \cdot \mathbf{x}_i + b) \geq 1$

One constraint per training point.

Note sign trick:
 $\mathbf{w} \cdot \mathbf{x}_i + b \geq 1$ (if $y_i = 1$)

$\mathbf{w} \cdot \mathbf{x}_i + b \leq -1$ (if $y_i = -1$)
 $(-1) \mathbf{w} \cdot \mathbf{x}_i - b \geq 1$

Finding the maximum margin line

- Solution: $\mathbf{w} = \sum_i \alpha_i y_i \mathbf{x}_i$

The diagram illustrates the components of the weight vector \mathbf{w} . It shows a sum of terms, where each term is a learned weight α_i multiplied by a support vector $y_i \mathbf{x}_i$. Two red arrows point from boxes labeled "Learned weight" and "Support vector" to the α_i and $y_i \mathbf{x}_i$ components respectively.

Learned weight Support vector

Finding the maximum margin line

- Solution: $\mathbf{w} = \sum_i \alpha_i y_i \mathbf{x}_i$
 $b = y_i - \mathbf{w} \cdot \mathbf{x}_i$ (for any support vector)
- Classification function:

$$\begin{aligned} f(\mathbf{x}) &= \text{sign}(\mathbf{w} \cdot \mathbf{x} + b) \\ &= \text{sign}\left(\sum_i \alpha_i y_i \mathbf{x}_i \cdot \mathbf{x} + b\right) \end{aligned}$$

If $f(\mathbf{x}) < 0$, classify as negative, otherwise classify as positive.

- Notice that it relies on an *inner product* between the test point \mathbf{x} and the support vectors \mathbf{x}_i
- (Solving the optimization problem also involves computing the *inner products* $\mathbf{x}_i \cdot \mathbf{x}_j$ between all pairs of training points)

Inner product

- The decision boundary for the SVM and its optimization depend on the inner product of two data points (vectors):

$$\mathbf{x}_i^T \mathbf{x}_j \quad f(x) = \text{sign}(\mathbf{w} \cdot \mathbf{x} + b) \\ = \text{sign}\left(\sum_i \alpha_i y_i \mathbf{x}_i^T \mathbf{x} + b\right)$$

- The inner product is equal

$$(\mathbf{x}_i^T \mathbf{x}) = \|\mathbf{x}_i\| * \|\mathbf{x}_i\| \cos \theta$$

If the angle in between them is 0 then: $(\mathbf{x}_i^T \mathbf{x}) = \|\mathbf{x}_i\| * \|\mathbf{x}_i\|$

If the angle between them is 90 then: $(\mathbf{x}_i^T \mathbf{x}) = 0$

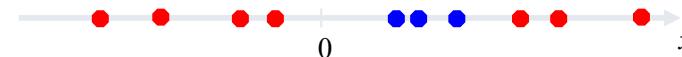
The inner product measures how similar the two vectors are

Nonlinear SVMs

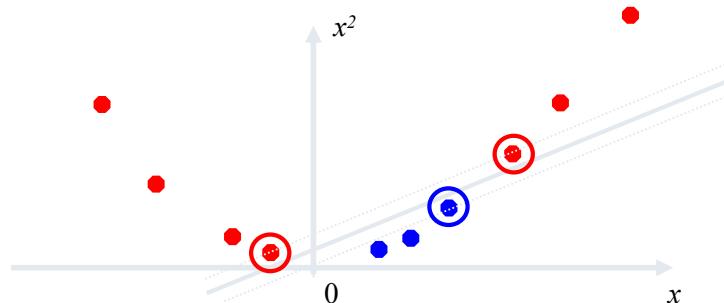
- Datasets that are linearly separable work out great:



- But what if the dataset is just too hard?



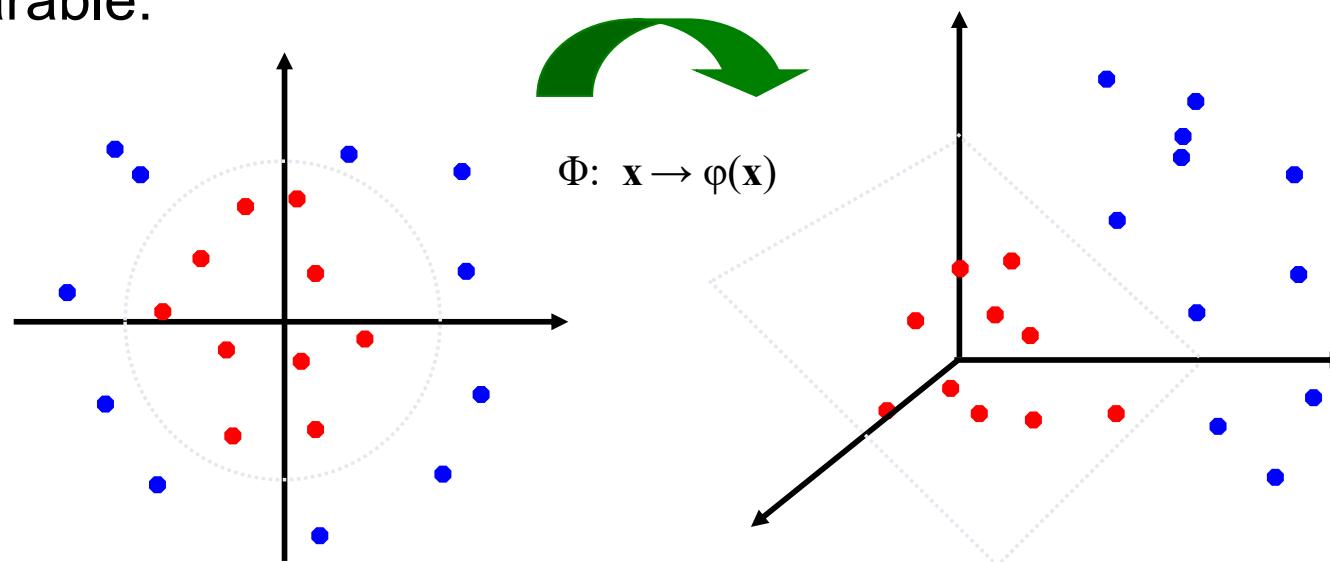
- We can map it to a higher-dimensional space:



Andrew Moore

Nonlinear SVMs

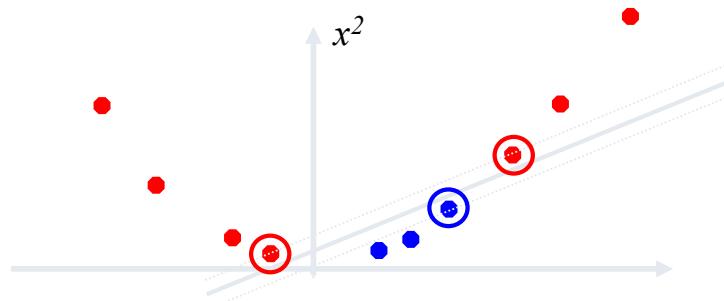
- **General idea:** the original input space can always be mapped to some higher-dimensional feature space where the training set is separable:



Nonlinear kernel: Example

- Consider the mapping

$$\varphi(x) = (x, x^2)$$



$$\varphi(x) \cdot \varphi(y) = (x, x^2) \cdot (y, y^2) = xy + x^2y^2$$

$$K(x, y) = xy + x^2y^2$$

The “Kernel Trick”

- The linear classifier relies on dot product between vectors $K(\mathbf{x}_i, \mathbf{x}_j) = \mathbf{x}_i \cdot \mathbf{x}_j$
- If every data point is mapped into high-dimensional space via some transformation $\Phi: \mathbf{x}_i \rightarrow \varphi(\mathbf{x}_i)$, the dot product becomes: $K(\mathbf{x}_i, \mathbf{x}_j) = \varphi(\mathbf{x}_i) \cdot \varphi(\mathbf{x}_j)$
- A *kernel function* is similarity function that corresponds to an inner product in some expanded feature space
- *The kernel trick*: instead of explicitly computing the lifting transformation $\varphi(\mathbf{x})$, define a kernel function K such that: $K(\mathbf{x}_i, \mathbf{x}_j) = \varphi(\mathbf{x}_i) \cdot \varphi(\mathbf{x}_j)$

Examples of kernel functions

- Linear:

$$K(x_i, x_j) = x_i^T x_j$$

- Polynomials of degree up to d :

$$K(x_i, x_j) = (x_i^T x_j + 1)^d$$

- Gaussian RBF:

$$K(x_i, x_j) = \exp\left(-\frac{\|x_i - x_j\|^2}{2\sigma^2}\right)$$

- Histogram intersection:

$$K(x_i, x_j) = \sum_k \min(x_i(k), x_j(k))$$

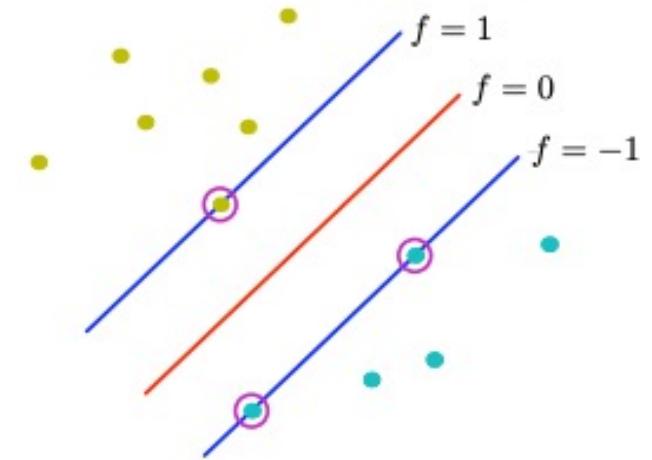
Hard-margin SVMs

$$\min_{\mathbf{w}} \frac{1}{2} \|\mathbf{w}\|^2$$

The \mathbf{w} that minimizes...

$\underbrace{}$

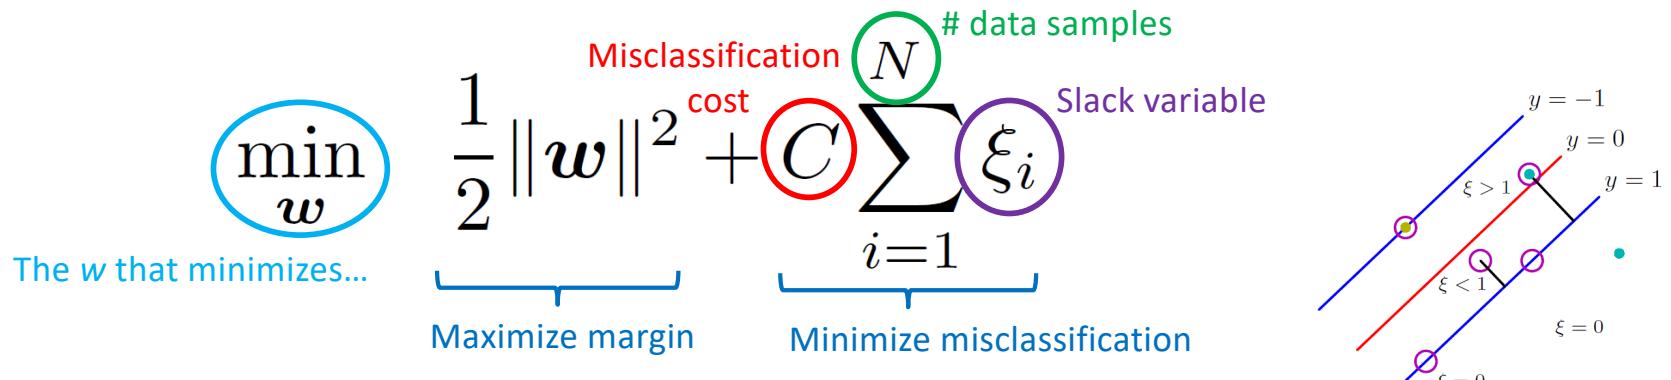
Maximize margin



subject to $y_i \mathbf{w}^T \mathbf{x}_i \geq 1 \quad ,$

$\forall i = 1, \dots, N$

Soft-margin SVMs



subject to

$$y_i w^T x_i \geq 1 - \xi_i,$$

$$\xi_i \geq 0, \quad \forall i = 1, \dots, N$$

Figure from Chris Bishop

Soft-margin SVMs

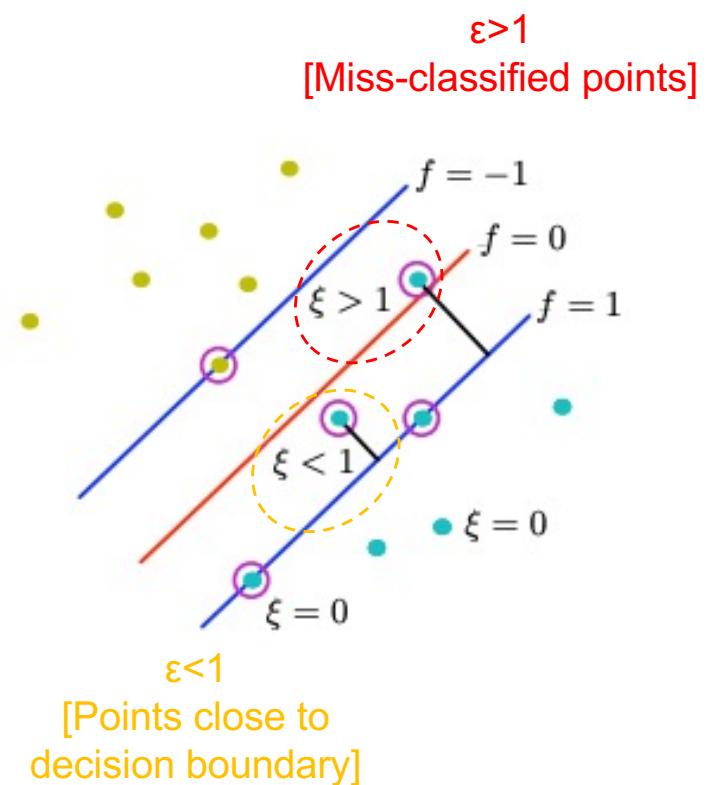
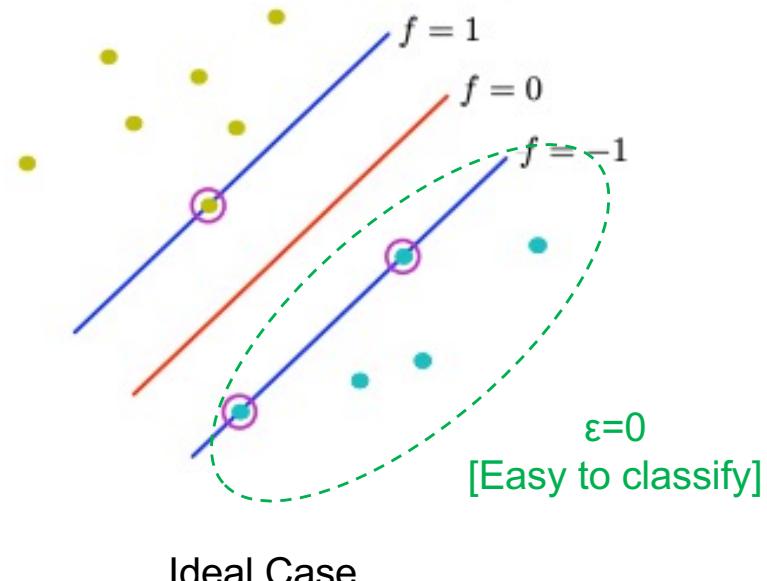
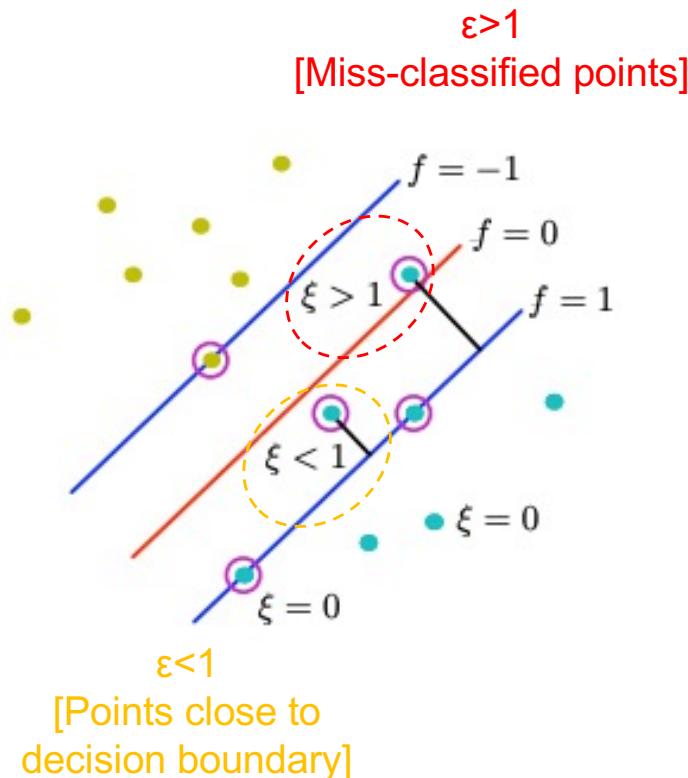


Figure from Chris Bishop

Soft-margin SVMs



Slack variables allow:

- Certain training points can be within the margin.
- We want these number of points as small as possible.

How do we minimize the second term in the optimization?

- A lot of examples with $\xi=0$ (easy correctly classified)
- Medium quantity of examples with $0<\xi<1$ (correct classified inside margin)
- Few examples with $\xi>1$ (misclassified examples)

Figure from Chris Bishop

What about multi-class SVMs?

- Unfortunately, there is no “definitive” multi-class SVM formulation
- In practice, we have to obtain a multi-class SVM by combining multiple two-class SVMs
- One vs. others/all
 - **Training**: learn an SVM for each class vs. the others
 - **Testing**: apply each SVM to the test example, and assign it to the class of the SVM that returns the highest decision value
- One vs. one
 - **Training**: learn an SVM for each pair of classes
 - **Testing**: each learned SVM “votes” for a class to assign to the test example

Multi-class problems

- One-vs-all (a.k.a. one-vs-others)
 - Train K classifiers
 - In each, pos = data from class i , neg = data from classes other than i
 - The class with the most confident prediction wins
 - Example:
 - You have 4 classes, train 4 classifiers
 - 1 vs others: score 3.5
 - 2 vs others: score 6.2
 - 3 vs others: score 1.4
 - 4 vs other: score 5.5
 - Final prediction: class 2

Multi-class problems

- One-vs-one (a.k.a. all-vs-all)
 - Train $K(K-1)/2$ binary classifiers (all pairs of classes)
 - They all vote for the label
 - Example:
 - You have 4 classes, then train 6 classifiers
 - 1 vs 2, 1 vs 3, 1 vs 4, 2 vs 3, 2 vs 4, 3 vs 4
 - Votes: 1, 1, 4, 2, 4, 4
 - Final prediction is class 4

Using SVMs

1. Select a kernel function.
2. Compute pairwise kernel values between labeled examples.
3. Use this “[kernel matrix](#)” to solve for SVM support vectors & alpha weights.
4. To classify a new example: compute kernel values between new input and support vectors, apply alpha weights, check sign of output.

Some SVM packages

- LIBSVM <http://www.csie.ntu.edu.tw/~cjlin/libsvm/>
- LIBLINEAR <https://www.csie.ntu.edu.tw/~cjlin/liblinear/>
- SVM Light <http://svmlight.joachims.org/>
- Scikit Learn <https://scikit-learn.org/stable/modules/svm.html>

Linear classifiers vs nearest neighbors

- Linear pros:
 - + Low-dimensional *parametric* representation
 - + Very fast at test time
- Linear cons:
 - Can be tricky to select best kernel function for a problem
 - Learning can take a very long time for large-scale problem
- NN pros:
 - + Works for any number of classes
 - + Decision boundaries not necessarily linear
 - + *Nonparametric* method
 - + Simple to implement
- NN cons:
 - Slow at test time (large search problem to find neighbors)
 - Storage of data
 - Especially need good distance function (but true for all classifiers)

Adapted from L. Lazebnik

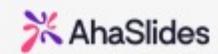
Lab 5: SVM

Duration: 30 min

Use JPEG, PNG and GIF files less than 15 MB [\[ahaslides\]](#)



To join, go to: ahaslides.com/FHBCJ



Please, run linear SVM on our face dataset and upload your resulted accuracy.

^ Get Feedback



Beyond Bags of Features: Spatial Pyramid Matching for Recognizing Natural Scene Categories

CVPR 2006

Winner of 2016 Longuet-Higgins Prize

Svetlana Lazebnik (slazebni@uiuc.edu)

Beckman Institute, University of Illinois at Urbana-Champaign

Cordelia Schmid

(cordelia.schmid@inrialpes.fr)

INRIA Rhône-Alpes, France

Jean Ponce ([ponce@di.ens.fr](mailto:pounce@di.ens.fr))

Ecole Normale Supérieure, France

Scene category dataset

Fei-Fei & Perona (2005), Oliva & Torralba (2001)

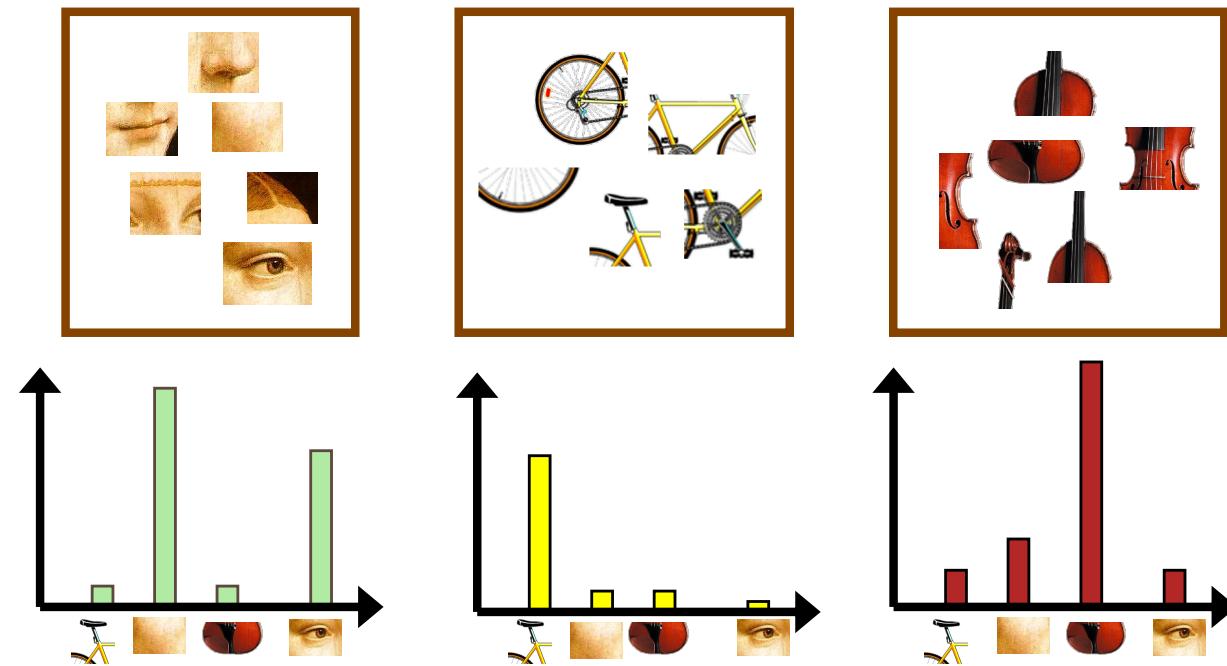
http://www-cvr.ai.uiuc.edu/ponce_grp/data



Slide credit: L. Lazebnik

Bag-of-words representation

1. Extract local features
2. Learn “visual vocabulary” using clustering
3. Quantize local features using visual vocabulary
4. Represent images by frequencies of “visual words”



Slide credit: L. Lazebnik

Image categorization with bag of words

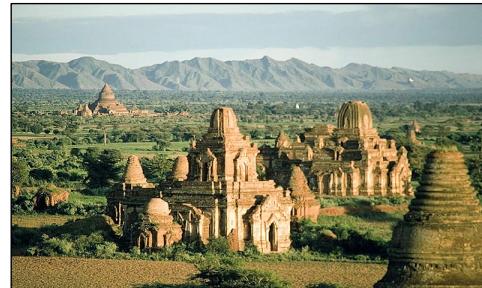
Training

1. Compute bag-of-words representation for training images
2. Train classifier on labeled examples using histogram values as features
3. Labels are the scene types (e.g. mountain vs field)

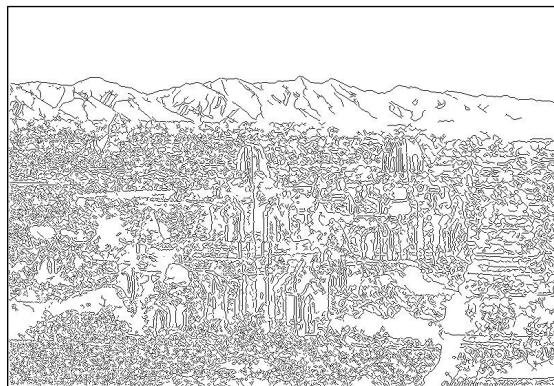
Testing

1. Extract keypoints / descriptors for test images
2. Quantize into visual words **using the clusters computed at training time**
3. Compute visual word histogram for test images
4. Compute labels on test images using classifier obtained at training time
5. **Evaluation only, do only once:** Measure accuracy of test predictions by comparing them to ground-truth test labels (obtained from humans)

Feature extraction (on which BOW is based)

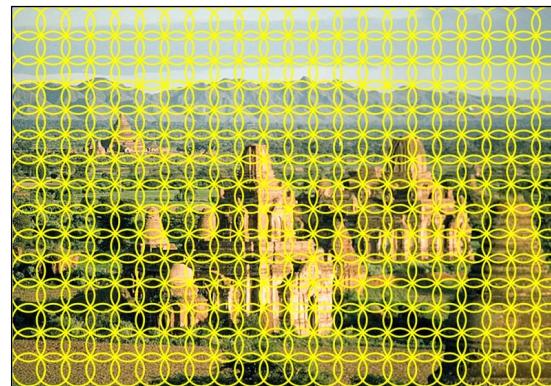


Weak features



Edge points at 2 scales and 8 orientations
(vocabulary size 16)

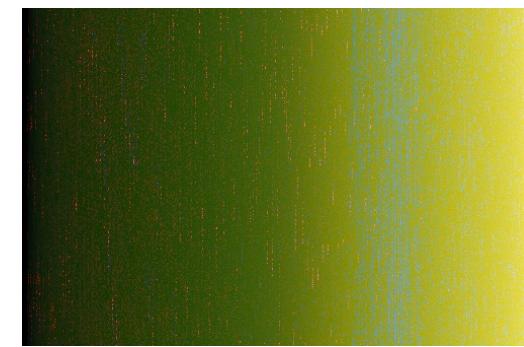
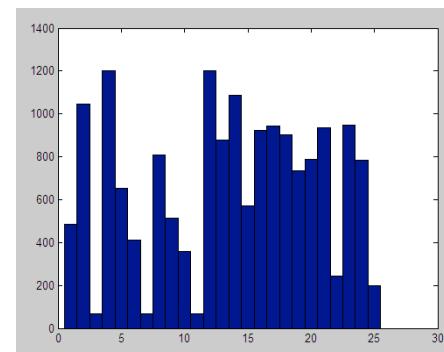
Strong features



SIFT descriptors of 16x16 patches sampled
on a regular grid, quantized to form visual
vocabulary (size 200, 400)

Slide credit: L. Lazebnik

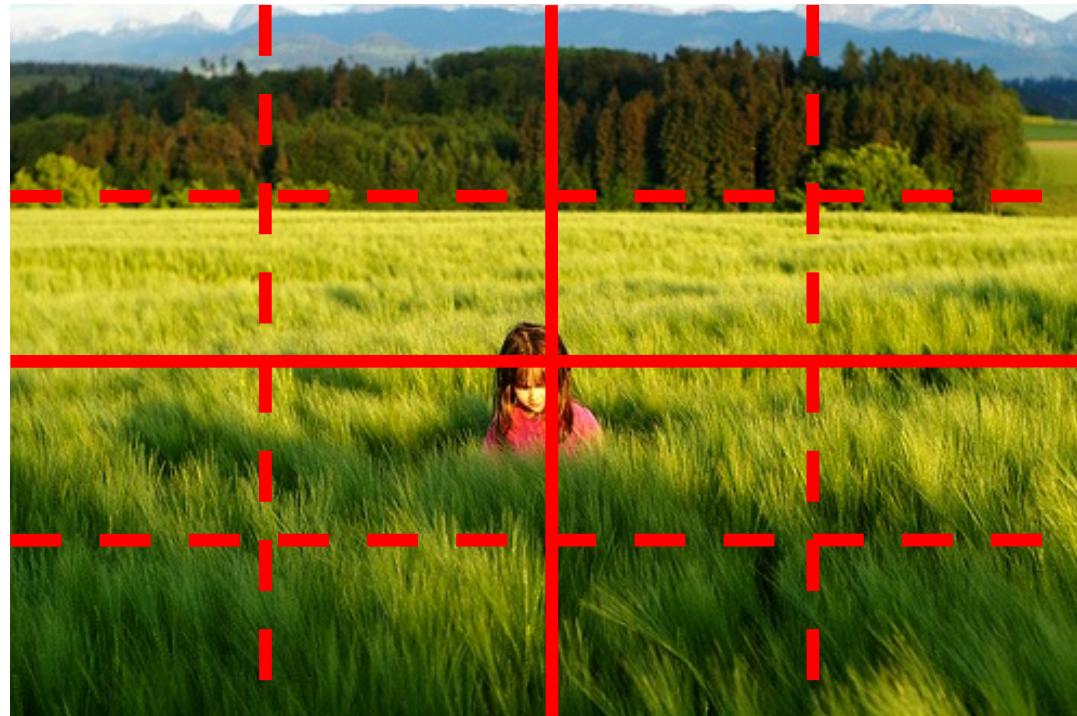
What about spatial layout?



All of these images have the same color histogram

Slide credit: D. Hoiem

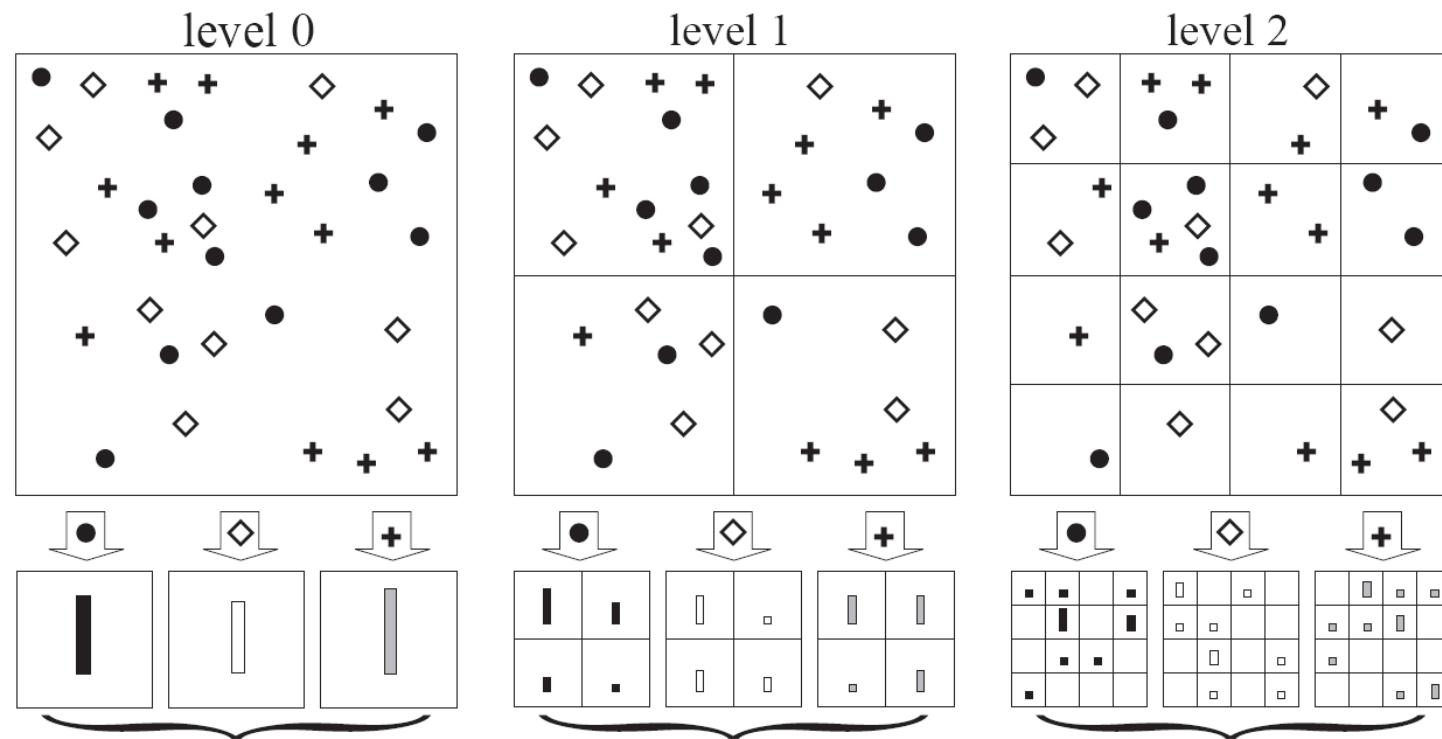
Spatial pyramid



Compute histogram in each spatial bin

Slide credit: D. Hoiem

Spatial pyramid

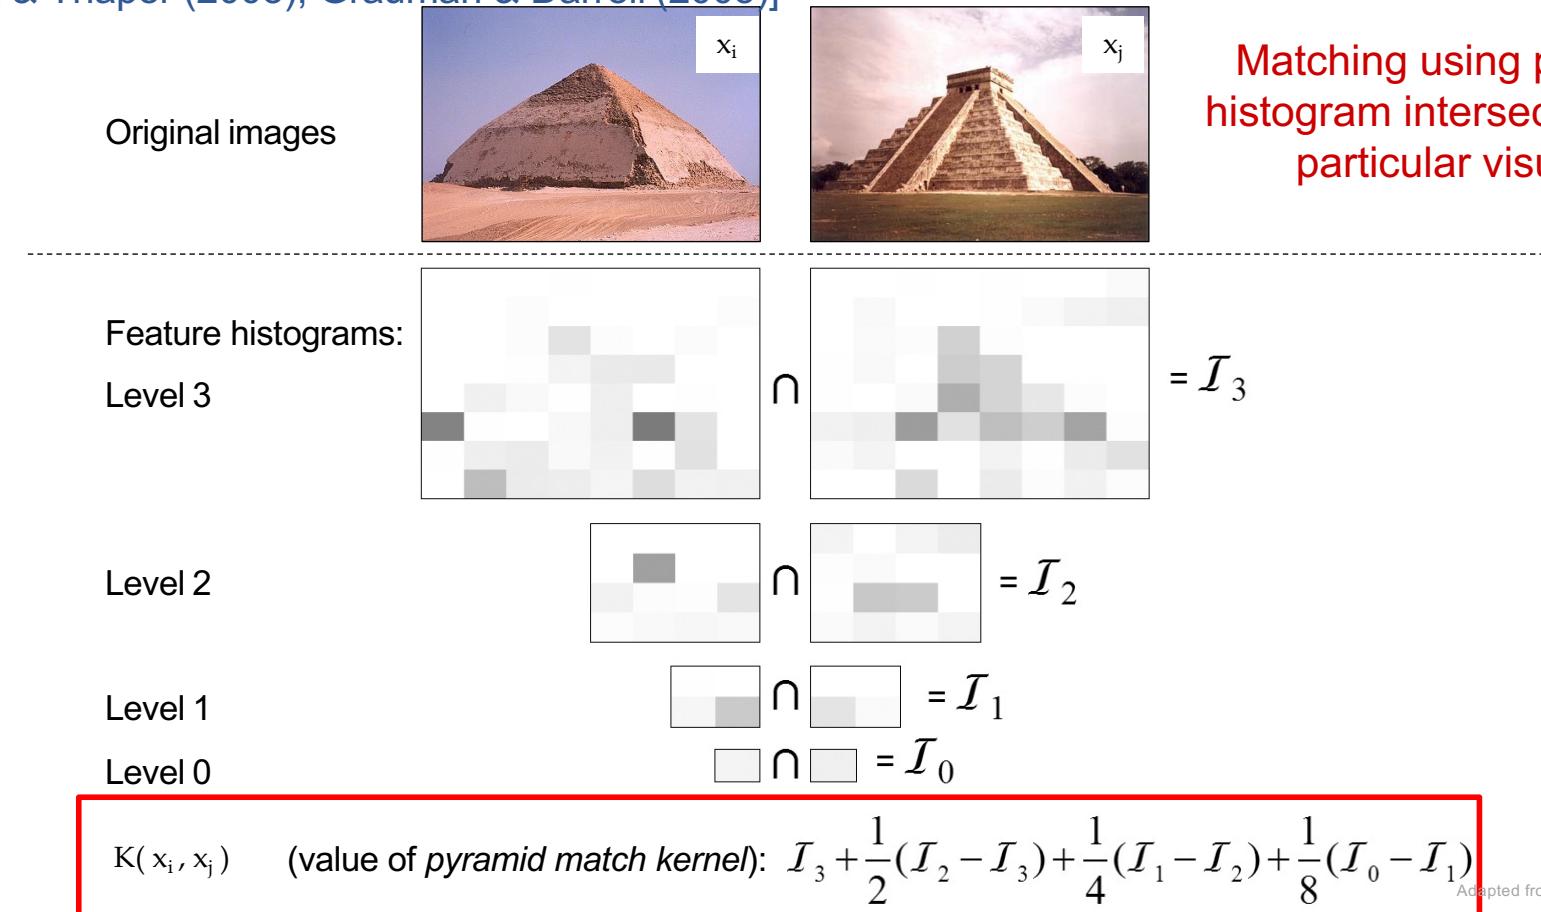


[Lazebnik et al. CVPR 2006]

Slide credit: D. Hoiem

Pyramid Matching

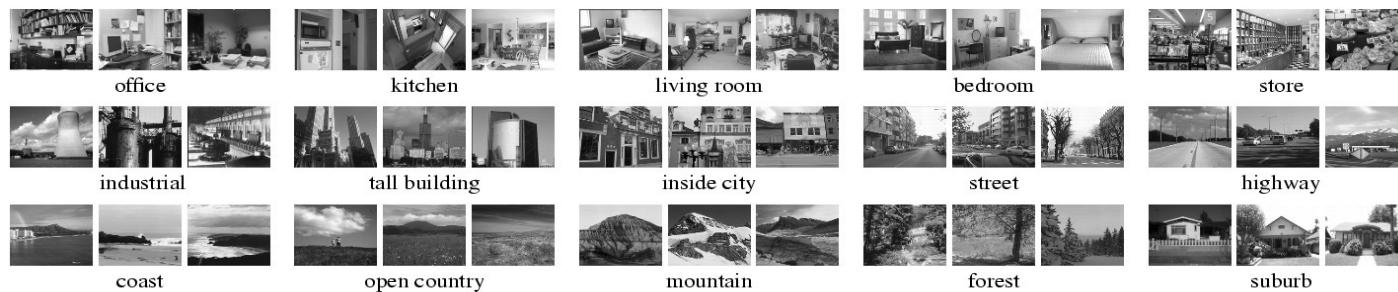
[Indyk & Thaper (2003), Grauman & Darrell (2005)]



Scene category dataset

Fei-Fei & Perona (2005), Oliva & Torralba (2001)

http://www-cvr.ai.uiuc.edu/ponce_grp/data



Multi-class classification results (100 training images per class)

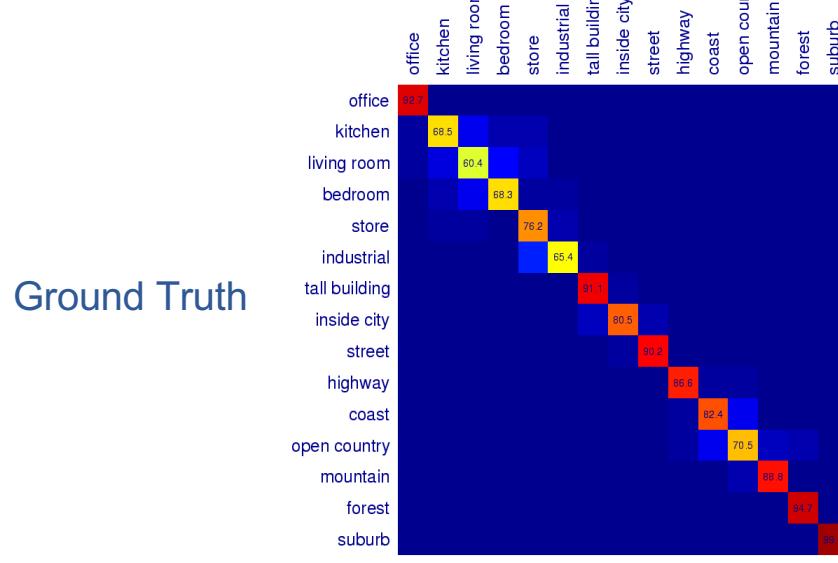
	Weak features (vocabulary size: 16)		Strong features (vocabulary size: 200)	
Level	Single-level	Pyramid	Single-level	Pyramid
0 (1×1)	45.3 ± 0.5		72.2 ± 0.6	
1 (2×2)	53.6 ± 0.3	56.2 ± 0.6	77.9 ± 0.6	79.0 ± 0.5
2 (4×4)	61.7 ± 0.6	64.7 ± 0.7	79.4 ± 0.3	81.1 ± 0.3
3 (8×8)	63.3 ± 0.8	66.8 ± 0.6	77.2 ± 0.4	80.7 ± 0.3

Fei-Fei & Perona: 65.2%

Slide credit: L. Lazebnik

Scene Category Confusion Matrix

Predictions



Ground Truth

Difficult indoor images



kitchen



living room



bedroom

Slide credit: L. Lazebnik

Caltech101 dataset

Fei-Fei et al. (2004)

http://www.vision.caltech.edu/Image_Datasets/Caltech101/Caltech101.html

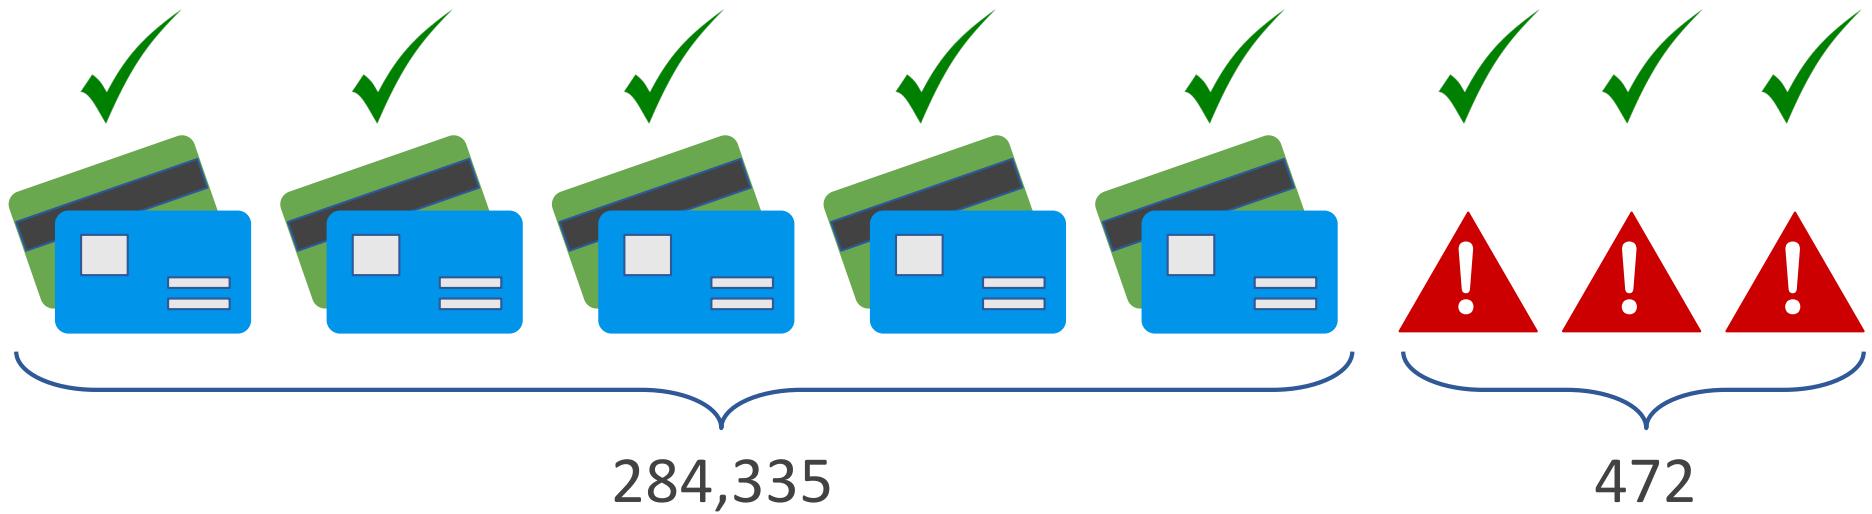


Multi-class classification results (30 training images per class)

		Weak features (16)		Strong features (200)	
Level		Single-level	Pyramid	Single-level	Pyramid
0		15.5 ± 0.9		41.2 ± 1.2	
1		31.4 ± 1.2	32.8 ± 1.3	55.9 ± 0.9	57.0 ± 0.8
2		47.2 ± 1.1	49.3 ± 1.4	63.6 ± 0.9	64.6 ± 0.8
3		52.2 ± 0.8	54.0 ± 1.1	60.3 ± 0.9	64.6 ± 0.7

Slide credit: L. Lazebnik

Evaluation Metrics: Credit Card Fraud

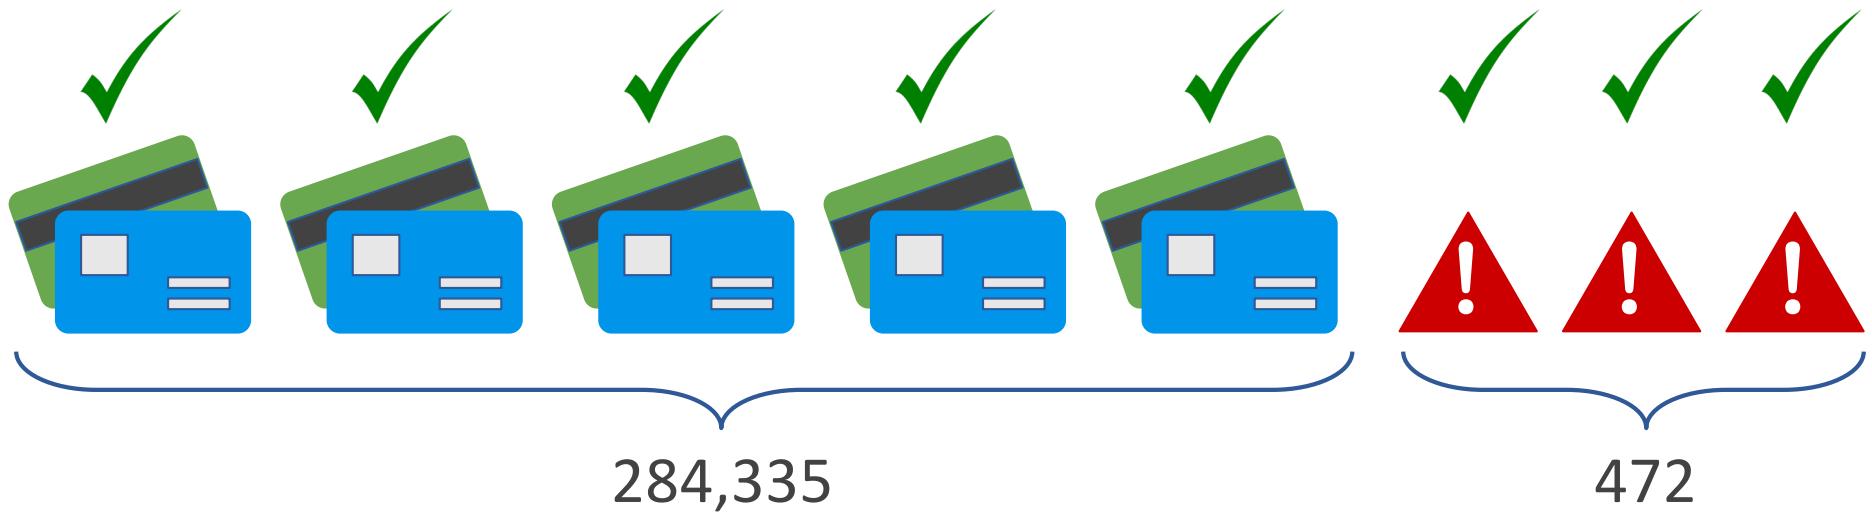


Model: All transactions are good.

$$\text{Correct} = \frac{284,335}{284,807} = 99.83\%$$

Slide Credit: Prof. Sandra Avila - UNICAMP

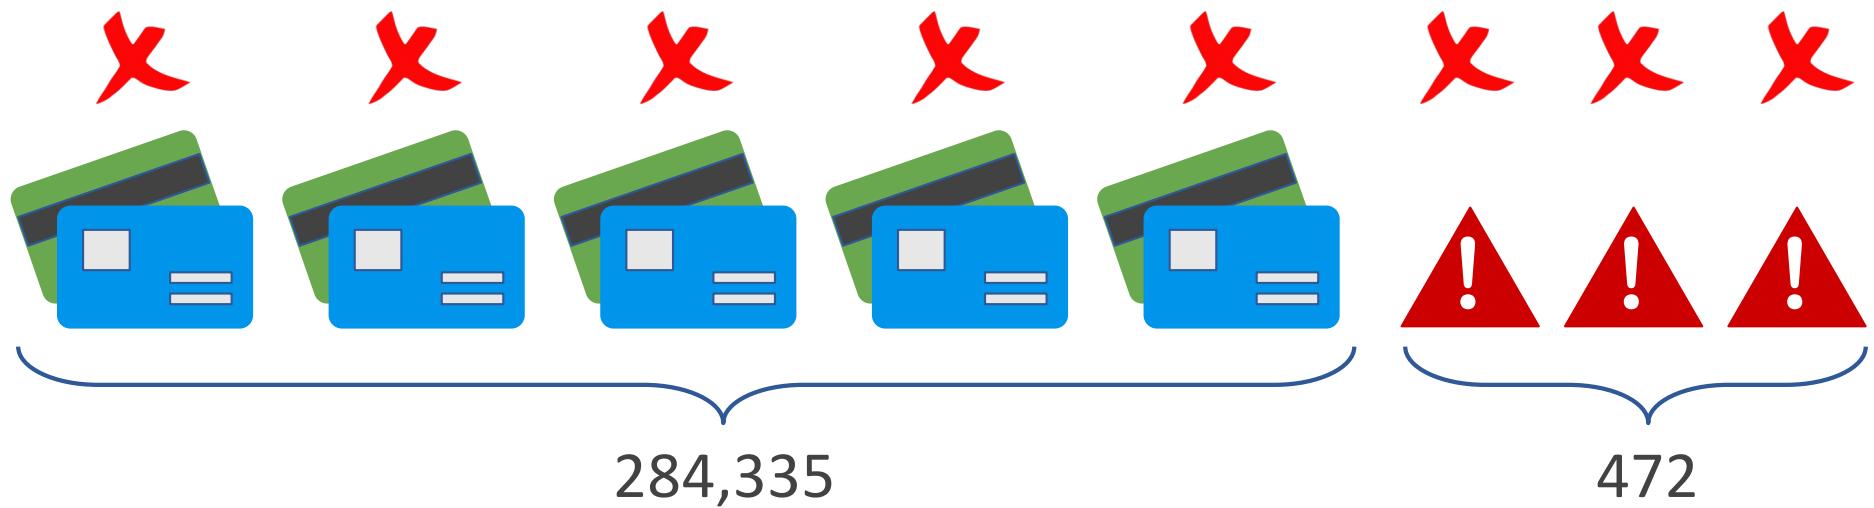
Evaluation Metrics: Credit Card Fraud



Model: All transactions are good.

Problem: I'm not catching any of the bad ones!

Evaluation Metrics: Credit Card Fraud



Model: All transactions are fraudulent.

Problem: I'm accidentally catching all the good ones!

Evaluation Metrics: Spam Classifier Model



Not Spam



Spam

Slide Credit: Prof. Sandra Avila - UNICAMP

Evaluation Metrics: Confusion Matrix Table

		Predictions	
		Sent to Spam Folder	Sent to Inbox
Annotations	Spam	True Positive 	False Negative 
	Not Spam	False Positive 	True Negative 

Slide Credit: Prof. Sandra Avila - UNICAMP

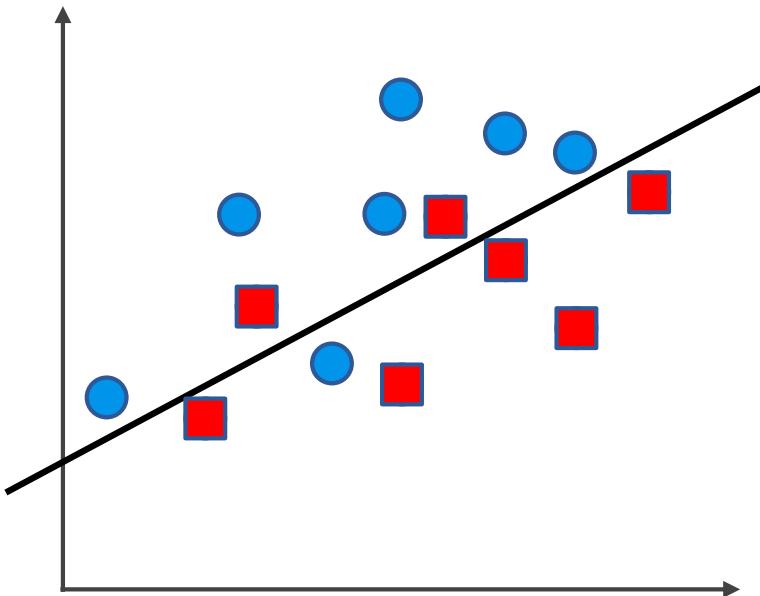
Slide Credit: Prof. Sandra Avila - UNICAMP

Evaluation Metrics: Confusion Matrix Table

		Folder	
		Spam Folder	Inbox
Email	Spam	100	170
	Not Spam	30	700

Slide Credit: Prof. Sandra Avila - UNICAMP

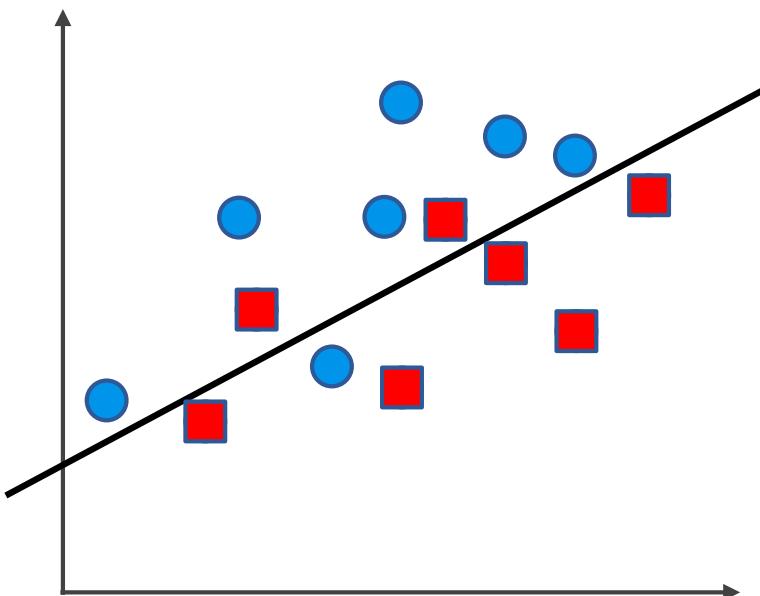
Evaluation Metrics: Confusion Matrix Table



		Prediction	
		Guessed Positive	Guessed Negative
Data	Positive		
	Negative		

Slide Credit: Prof. Sandra Avila - UNICAMP

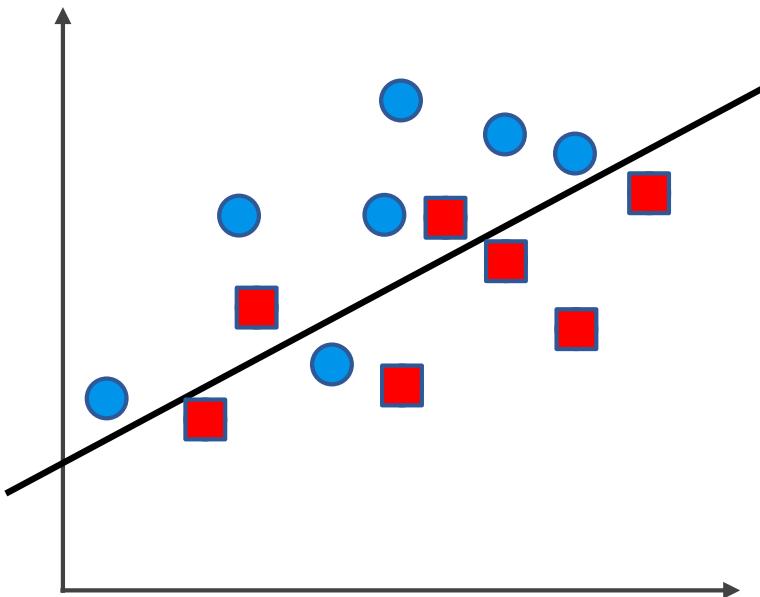
Evaluation Metrics: Confusion Matrix Table



		Prediction	
		Guessed Positive	Guessed Negative
Data	Positive	6 True positives	
	Negative		

Slide Credit: Prof. Sandra Avila - UNICAMP

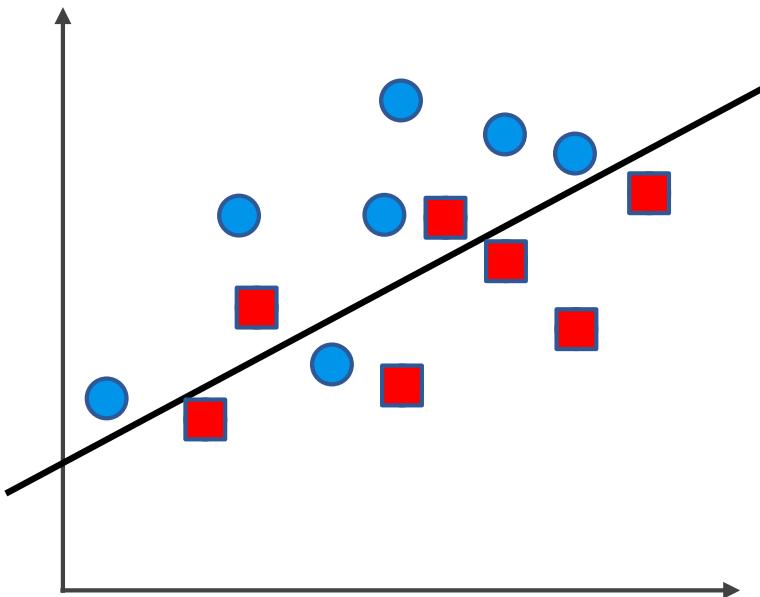
Evaluation Metrics: Confusion Matrix Table



		Prediction	
		Guessed Positive	Guessed Negative
Data	Positive	6 True positives	
	Negative		5 True negatives

Slide Credit: Prof. Sandra Avila - UNICAMP

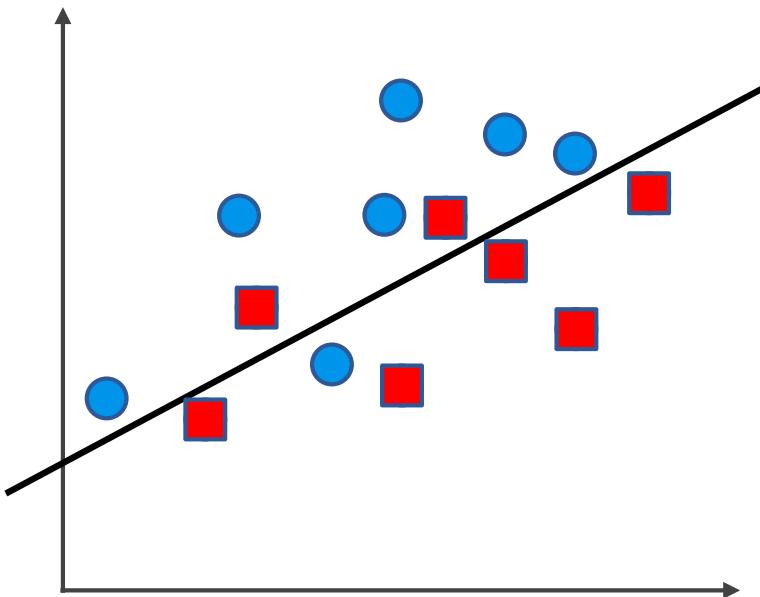
Evaluation Metrics: Confusion Matrix Table



		Prediction	
		Guessed Positive	Guessed Negative
Data	Positive	6 True positives	1 False negative
	Negative		5 True negatives

Slide Credit: Prof. Sandra Avila - UNICAMP

Evaluation Metrics: Confusion Matrix Table



		Prediction	
		Guessed Positive	Guessed Negative
Data	Positive	6 True positives	1 False negative
	Negative	2 False positives	5 True negatives

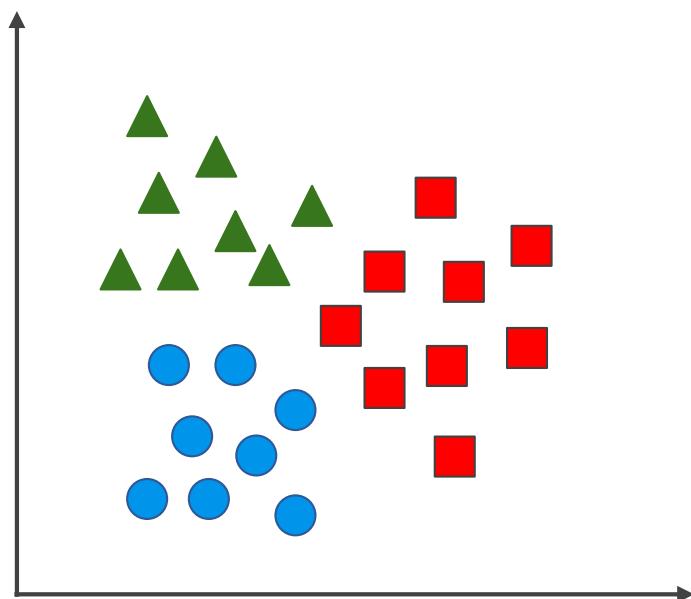
Slide Credit: Prof. Sandra Avila - UNICAMP

Evaluation Metrics: Confusion Matrix Table (n classes)

Class 1: ▲

Class 2: ■

Class 3: ●



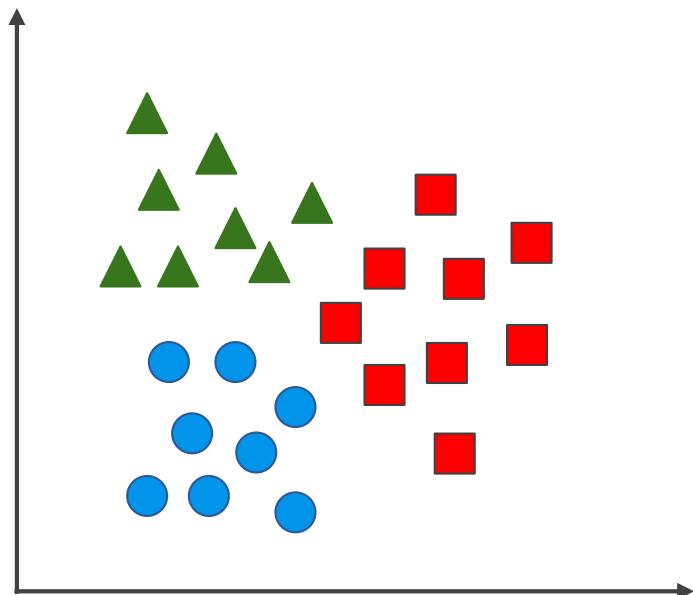
Slide Credit: Prof. Sandra Avila - UNICAMP

Evaluation Metrics: Confusion Matrix Table (n classes)

Class 1: ▲

Class 2: ■

Class 3: ●



True Class	Predicted Class		
	Guessed Class 1	Guessed Class 2	Guessed Class 3
Class 1			
Class 2			
Class 3			

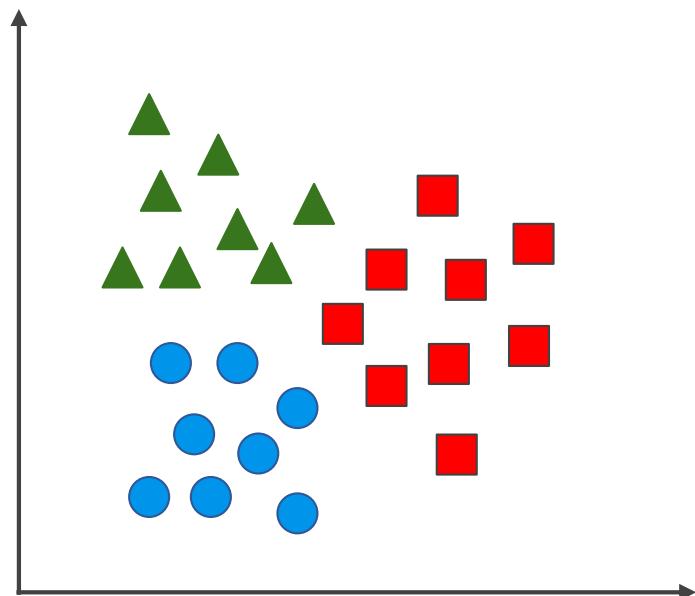
Slide Credit: Prof. Sandra Avila - UNICAMP

Evaluation Metrics: Confusion Matrix Table (n classes)

Class 1: ▲

Class 2: ■

Class 3: ●



		Predicted Class		
		Guessed Class 1	Guessed Class 2	Guessed Class 3
True Class	Class 1	5	2	1
	Class 2	3	6	0
	Class 3	0	1	7

Slide Credit: Prof. Sandra Avila - UNICAMP

Evaluation Metrics: Accuracy

		Diagnosis	
		Diagnosed Sick	Diagnosed Healthy
Patients	Sick	1,000	200
	Healthy	800	8,000

Slide Credit: Prof. Sandra Avila - UNICAMP

Evaluation Metrics: Accuracy

		Diagnosis	
		Diagnosed Sick	Diagnosed Healthy
Patients	Sick	1,000	200
	Healthy	800	8,000

Accuracy:

Out of all the **patients**, how many did we classify correctly?

Slide Credit: Prof. Sandra Avila - UNICAMP

Evaluation Metrics: Accuracy

		Diagnosis	
		Diagnosed Sick	Diagnosed Healthy
Patients	Sick	1,000	200
	Healthy	800	8,000

Accuracy:

Out of all the **patients**, how many did we classify correctly?

Accuracy =

$$\frac{1,000 + 8,000}{10,000}$$

Evaluation Metrics: Accuracy

		Diagnosis	
		Diagnosed Sick	Diagnosed Healthy
Patients	Sick	1,000	200
	Healthy	800	8,000

Accuracy:

Out of all the **patients**, how many did we classify correctly?

Accuracy =

$$\frac{1,000 + 8,000}{10,000} = 90\%$$

Evaluation Metrics: Accuracy

Email	Folder	
	Spam Folder	Inbox
Spam	100	170
Not Spam	30	700

Accuracy:

Out of all the **emails**, how many did we classify correctly?

Evaluation Metrics: Accuracy

Email	Folder	
	Spam Folder	Inbox
Spam	100	170
Not Spam	30	700

Accuracy:

Out of all the **emails**, how many did we classify correctly?

Accuracy =

$$\frac{100 + 700}{1,000} = 80\%$$

Slide Credit: Prof. Sandra Avila - UNICAMP

Summary

- Visual Classification
- Models
 - K-Nearest Neighbor
 - Support Vector Machines (SVM)
- Spatial Pyramid Feature Extractor
- Evaluation Metrics

